





The brain is dynamic




Neural plasticity


Changing the synapse strength



Changing the synapse strength




Perceptrons



Frank
Rosenblatt

Multilayer perceptrons



Evolving neural networks through augmenting topologies

[KO Stanley, R Miikkulainen](#) - *Evolutionary computation*, 2002 - MIT Press



An important question in neuroevolution is how to gain an advantage from evolving neural network topologies along with weights. We present a method, NeuroEvolution of Augmenting Topologies (NEAT), which outperforms the best fixed-topology method on a challenging benchmark reinforcement learning task. We claim that the increased efficiency is due to (1) employing a principled method of crossover of different topologies, (2) protecting structural innovation using speciation, and (3) incrementally growing from minimal structure ...

☆ 99 Zitiert von: 3254 Ähnliche Artikel Alle 23 Versionen


Genome (Genotype)

Node Genes	Node 1	Node 2	Node 3	Node 4	Node 5			
	Sensor	Sensor	Sensor	Output	Hidden			
Connect. Genes	In 1	In 2	In 3	In 2	In 5	In 1	In 4	
	Out 4	Out 4	Out 4	Out 5	Out 4	Out 5	Out 5	
	Weight 0.7	Weight-0.5	Weight 0.5	Weight 0.2	Weight 0.4	Weight 0.6	Weight 0.6	
	Enabled	DISABLED	Enabled	Enabled	Enabled	Enabled	Enabled	
	Innov 1	Innov 2	Innov 3	Innov 4	Innov 5	Innov 6	Innov 11	

Network (Phenotype)



$$\begin{array}{r} [A,B,C] \\ \times [C,B,A] \\ \hline \end{array}$$
 Crossovers: [A,B,A] [C,B,C]
 (both are missing information)



„**MarI/O** is a program made of neural networks and genetic algorithms that kicks butt at Super Mario World.”


© Seth Bling, YouTube

<https://www.youtube.com/watch?v=qv6UVOQ0F44>

How to change weights?


Hebb's rule:

„Neurons that fire together, wire together.“




$$w_{ij} = x_i x_j$$

x_i	x_j	w_{ij}
0	0	0
1	0	0
0	1	0
1	1	1




Structural
plasticity



Synaptic
plasticity


Spike timing dependent plasticity (STDP)

Hebb's rule – cntd.



< 20 ms

Synaptic $w \uparrow \rightarrow$ LTP




< 20 ms

Synaptic $w \downarrow \rightarrow$ LTD

Spike-timing-dependent plasticity rewards synchrony rather than causality


Margarita Anisimova¹, Bas van Bommel¹, Marina Mikhaylova^{1,2}, J. Simon Wiegert, Thomas G.

Oertner^{1*}, Christine E. Gee^{1,3*}



<https://doi.org/10.1101/863365>

“Our results confirm that neurons wire together if they fire together, but suggest that synaptic depression after anti-causal activation (tLTD) is a transient phenomenon.”



$$v = w \cdot u^T$$

The basic Hebb's rule:

$$\tau_w \frac{dw}{dt} = v \cdot u$$

$$\tau_w \frac{dw}{dt} = w \cdot \underbrace{u \cdot u^T}_{\text{Correlation matrix}}$$

$$\tau_w \frac{dw}{dt} = Q \cdot w^T$$

Correlation-based plasticity rule

$$w \rightarrow w + \epsilon Q \cdot w^T \quad \epsilon = \frac{1}{\tau_w}$$

Correlation-based plasticity rule

$$\tau_w \frac{dw}{dt} = \mathbf{Q} \cdot \mathbf{w}^T$$

→ Basic Hebb only allows LTP

Postsynaptic
LTD/LTP switch

$$\tau_w \frac{dw}{dt} = (v - \theta_v) \cdot \mathbf{u}$$

$$\tau_w \frac{dw}{dt} = v \cdot (\mathbf{u} - \theta_u)$$

Presynaptic
LTD/LTP switch

Covariance matrix

$$v = \mathbf{w} \cdot \mathbf{u}^T \quad \Rightarrow \quad \tau_w \frac{dw}{dt} = \mathbf{w} \cdot \mathbf{u} \cdot (\mathbf{u} - \theta_u)^T \quad \Rightarrow \quad \tau_w \frac{dw}{dt} = \mathbf{C} \cdot \mathbf{w}^T$$

Covariance-based plasticity rule

Hebbian learning suffers from instability

$$\tau_w \frac{dw}{dt} = v \cdot \mathbf{u} \cdot (v - \theta_v) \quad \text{If constant} \rightarrow \text{unstable}$$

→ Threshold of postsynaptic activity that determines if synapse is strengthened or weakened.

Adapt threshold θ_v :

$$\tau_\theta \frac{d\theta_v}{dt} = v^2 - \theta_v$$

$$\tau_{\theta} \frac{d\theta_v}{dt} = v^2 - \theta_v \quad \rightarrow \text{Stabilize weights through postsynaptic activity}$$

Can we use penalty terms directly on the weight vector?

$$\tau_w \frac{d\mathbf{w}}{dt} = v \cdot \mathbf{u} - \frac{v(\mathbf{n} \cdot \mathbf{u}^T)\mathbf{n}}{N_u}$$

Normalize by subtracting the same quantity

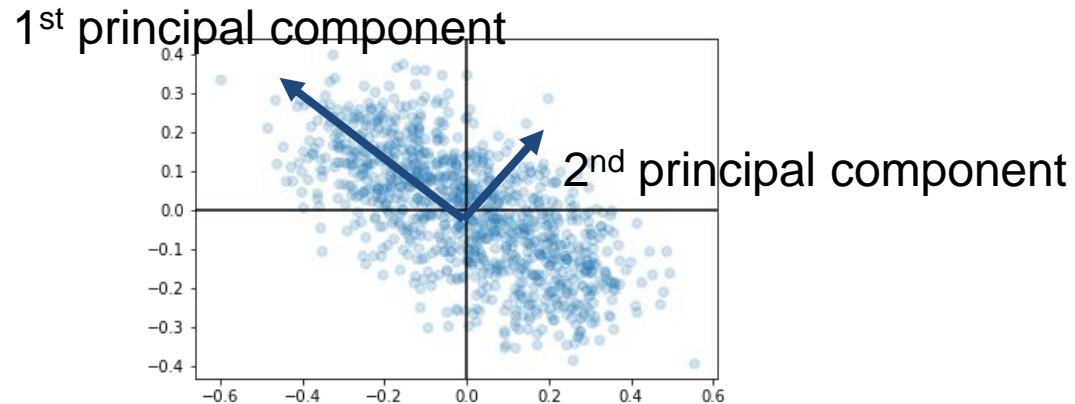


$$\tau_w \frac{d\mathbf{w}}{dt} = v \cdot \mathbf{u} - \alpha \cdot v^2 \cdot \mathbf{w} \quad \alpha > 0$$

$$\tau_w \frac{d\mathbf{w}}{dt} = v \cdot (\mathbf{u} - \alpha \cdot v \cdot \mathbf{w})$$

Unsupervised learning

PRINCIPAL COMPONENT ANALYSIS (PCA)




Simulation

Steps:

1. Data generation
2. Variable initialization
3. Iterate through data
 1. Compute pre-synaptic input
 2. Compute post-synaptic activation
 3. Compute Δw and update
4. Plot result

```
1 # That you see the same what I see
2 np.random.seed(42)
3 N = 1000
4
5 # Generate random data
6 x = np.linspace(-.3, .3, N)
7 np.random.shuffle(x)
8 y = -.7 * x
9
10 x += np.random.randn(x.size) / 10
11 y += np.random.randn(y.size) / 10
12
```




```
1 # Initialize some random weights
2 w = np.array([0.1, 0.4])
3 # Training iterations
4 N = 1000
5 eta = 0.1
6 ws = []
7
```


```
8 # Training
9 for i in range(N):
10     rPre = np.asarray([x[ix], y[ix]])
11     rPost = w @ rPre
12     w = w + eta * rPost * (rPre - rPost * w)
```

Simulation


Steps:

1. Data generation
2. Variable initialization
3. Iterate through data
 1. Compute pre-synaptic input
 2. Compute post-synaptic activation
 3. Compute Δw and update
4. Plot result






Perceptron



Unsupervised learning → no target imposed
Supervised learning → target imposed



$$\hat{y} = \sum_i w_i \cdot x_i$$

Target y : 2

Prediction \hat{y} : 1

$$\begin{aligned} \text{Error} &= \text{„Target – Prediction“} = \\ &= \frac{1}{2} (y - \hat{y})^2 \end{aligned}$$

$$\frac{dE}{dw_i} E = \frac{dE}{dw_i} \frac{1}{2} (y - \hat{y})^2 = -(y - \hat{y}) \frac{d\hat{y}}{dw_i} \sum_i w_i \cdot x_i = -(y - \hat{y}) \cdot x_i$$

Delta rule

$$\Delta w_i = \alpha \cdot (y - \hat{y}) \cdot x_i$$



Steps:

1. Data generation
2. Variable initialization
3. Iterating
 1. Compute prediction
 2. Update weights
4. Plotting


```
1 np.random.seed(42)
2
3 N = 100
4
5 xs = []
6
7 for _ in range(N):
8     x0 = np.random.randn()-1
9     x1 = np.random.randn()-1
10    xs.append((x0, x1))
11
12    x0 = np.random.randn()+2
13    x1 = np.random.randn()+1
14    xs.append((x0, x1))
15
16 xs = np.asarray(xs)
```

```
1 w = np.random.randn(2)
2 eta = 0.01
```

```
5 for i in range(2*N):
6     # Determine class/target
7     t = -1 if i % 2 == 0 else 1
8
9     # Compute prediction
10    pred = w @ xs[i]
11
12    # Update weights
13    w = w + eta * (t-pred) * xs[i]
14
```



Simulation over time




The brain



How does it emerge?

Back in the days...







© Pauline Breijer




What is happening there?





A





What is happening there?






Signaling pathways define neural development









Hox genes determines motor neurons




Positional development is highly conserved



Hoxa, chromosome 6



Hoxb, chromosome 11



Hoxc, chromosome 15



Hoxd, chromosome 2



Mammals have very similar development

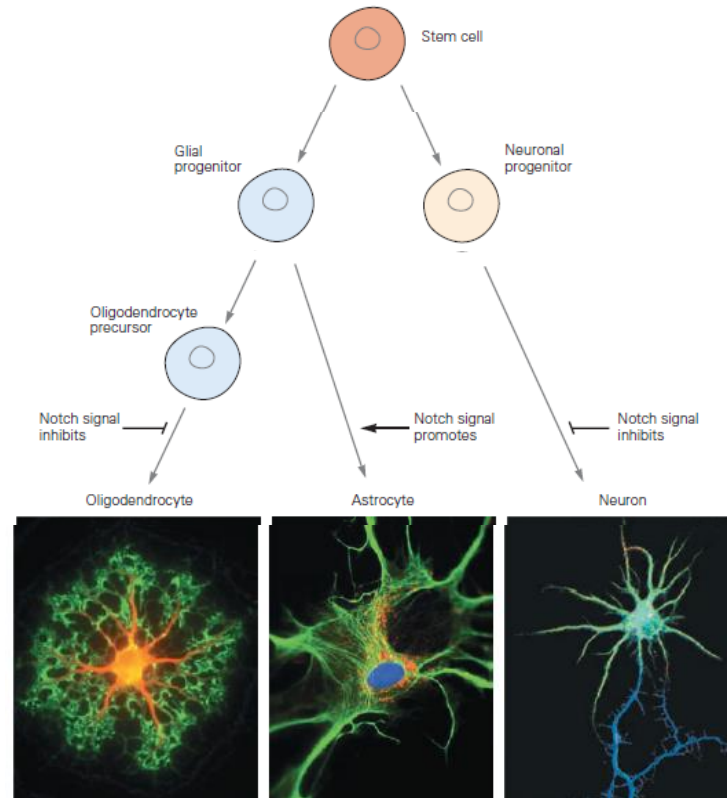


Human




Pig

Brain cells have a common ancestor





Neuron maturation




Pruning over lifetime

Experience Shapes Brain Architecture by Over-Production Followed by Pruning


Center on the Developing Child  HARVARD UNIVERSITY



Source: Shonkoff, J. P. (2008) **



From Cowan, 1979. © 1979 by Scientific American, Inc. All rights reserved.




From Huttenlocher and Dabholkar, 1997. Reprinted with permission of John Wiley & Sons Inc.

Optimal Brain Damage


Yann Le Cun, John S. Denker and Sara A. Solla
AT&T Bell Laboratories, Holmdel, N. J. 07733

ABSTRACT


We have used information-theoretic ideas to derive a class of practical and ~~nearly optimal schemes for adapting the size of a neural network~~. By removing unimportant weights from a network, several improvements can be expected: better generalization, fewer training examples required, and improved speed of learning and/or classification. The basic idea is to use second-derivative information to make a tradeoff between network complexity and training set error. Experiments confirm the usefulness of the methods on a real-world application.



Pruning synapses:
making network sparse
Pruning neurons:
Making network dense

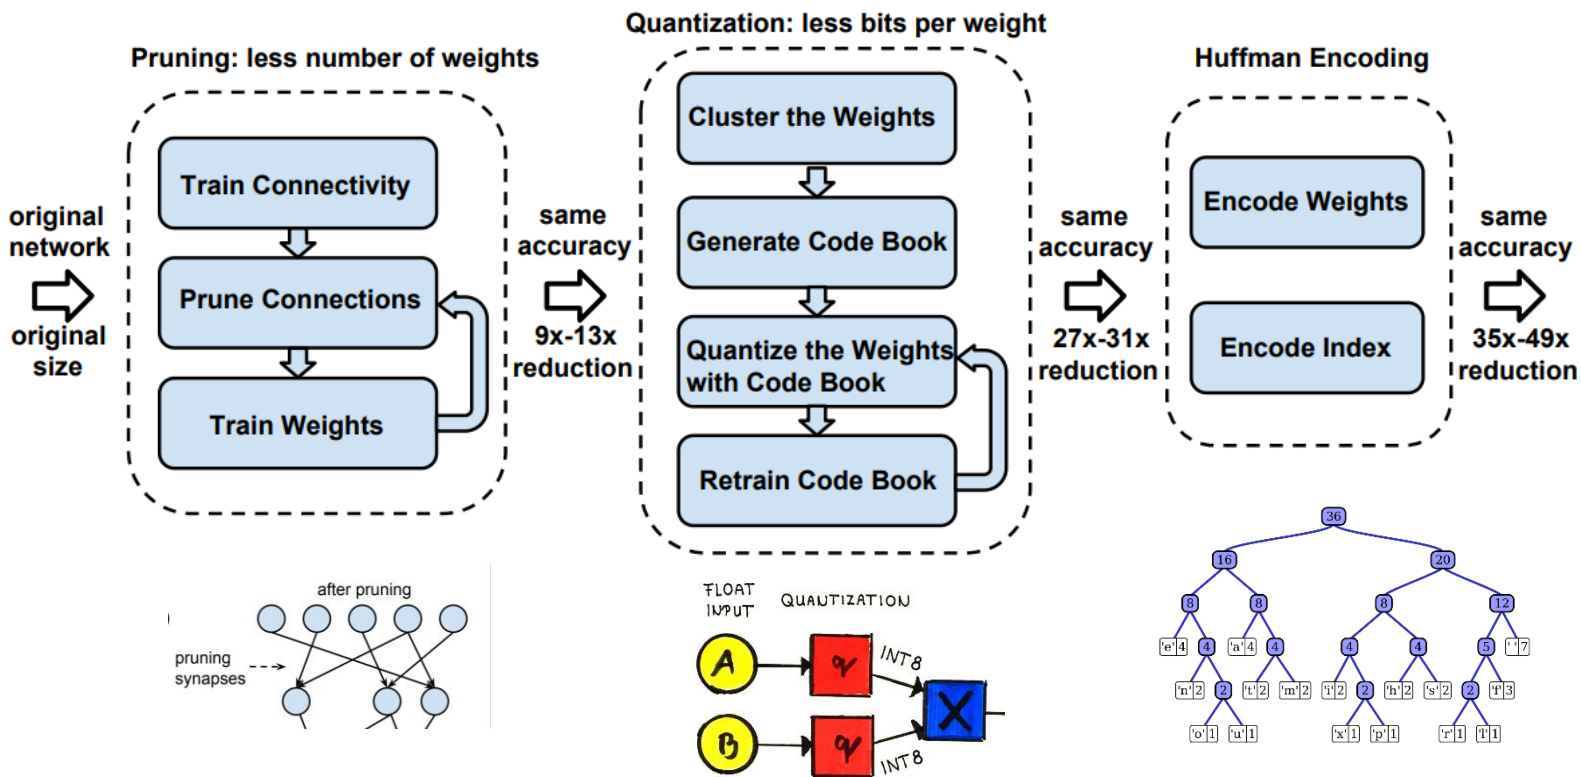


Layer	Weights	FLOP	Act%	Weights%	FLOP%
conv1	35K	211M	88%	84%	84%
conv2	307K	448M	52%	38%	33%
conv3	885K	299M	37%	35%	18%
conv4	663K	224M	40%	37%	14%
conv5	442K	150M	34%	37%	14%
fc1	38M	75M	36%	9%	3%
fc2	17M	34M	40%	9%	3%
fc3	4M	8M	100%	25%	10%
Total	61M	1.5B	54%	11%	30%

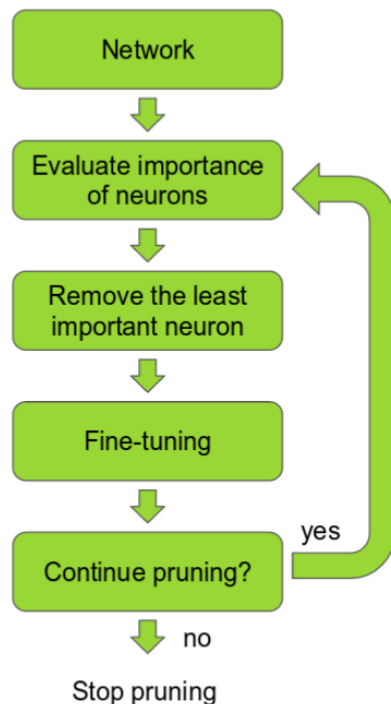


Layer	Weights	FLOP	Act%	Weights%	FLOP%
conv1_1	2K	0.2B	53%	58%	58%
conv1_2	37K	3.7B	89%	22%	12%
conv2_1	74K	1.8B	80%	34%	30%
conv2_2	148K	3.7B	81%	36%	29%
conv3_1	295K	1.8B	68%	53%	43%
conv3_2	590K	3.7B	70%	24%	16%
conv3_3	590K	3.7B	64%	42%	29%
conv4_1	1M	1.8B	51%	32%	21%
conv4_2	2M	3.7B	45%	27%	14%
conv4_3	2M	3.7B	34%	34%	15%
conv5_1	2M	925M	32%	35%	12%
conv5_2	2M	925M	29%	29%	9%
conv5_3	2M	925M	19%	36%	11%
fc6	103M	206M	38%	4%	1%
fc7	17M	34M	42%	4%	2%
fc8	4M	8M	100%	23%	9%
total	138M	30.9B	64%	7.5%	21%

Pruning in deep neural networks



Han et al., Deep Compression ICLR 2016



2.1 ORACLE PRUNING

Minimizing the difference in accuracy between the full and pruned models depends on the criterion for identifying the “least important” parameters, called *saliency*, at each step. The best criterion would be an exact empirical evaluation of each parameter, which we denote the *oracle* criterion, accomplished by ablating each non-zero parameter $w \in \mathcal{W}'$ in turn and recording the cost’s difference.

To compute the oracle, we evaluate the change in loss caused by removing each individual feature map from the fine-tuned VGG-16 network. (See Appendix A.3 for additional analysis.) We rank




Figure 4: Pruning of feature maps in VGG-16 fine-tuned on the Birds-200 dataset.






FIGURE 6. Relation between cell-planes and hypercolumns within a layer.

Deep Belief Networks



Restricted
Boltzmann
Machines

First: unsupervised → pre-training
Second: supervised → classification



Geoffrey Hinton, © WIRED




Table 9
Classification results.

Dataset	Net. size	DBN		CA	
		Acc. (%)	Connec. (%)	Acc. (%)	Connec. (%)
MNIST	N1	38.8	88.1	96.2	89.8
	N2	84.1	91.8	97.1	85.2
	N3	84.2	98.8	98.5	79.1
	N4	89.0	47.2	98.5	58.7
	N5	88.3	54.9	99.2	43.5
	N6	89.4	53.9	99.8	37.6
	N7	86.5	40.8	99.8	28.5



Neurodevelopment happens during our lifetime

