

# Lecture 3, Part 1

# Image Filtering (Linear Filters)

Computer Vision  
Summer Semester 2022

Prof. Bernhard Egger, Prof. Tim Weyrich, Prof. Andreas Maier

# Image Filtering

- Linear Filters -- Mean, Convolution, Gaussian
- Filter separability, etc

# Image filtering

- Image filtering:

- Compute function of local neighborhood at each position

$$h[m, n] = \sum_{k,l} f[k, l] I[m + k, n + l]$$

- Really important!
  - Enhance images → Denoise, resize, increase contrast, etc.
  - Extract information from images → Texture, edges, distinctive points, etc.
  - Detect patterns → Template matching

# Key properties of linear filters

- **Linearity:**

$$\text{imfilter}(I, f_1 + f_2) = \text{imfilter}(I, f_1) + \text{imfilter}(I, f_2)$$

- **Shift invariance:**

Same behavior regardless of pixel location

$$\text{imfilter}(I, \text{shift}(f)) = \text{shift}(\text{imfilter}(I, f))$$

*Any linear, shift-invariant operator can be represented as a convolution.*

# Correlation and Convolution

- 2d correlation 
$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$
  
`h=filter2D(f, I) / h=imfilter(I, f)` **[matlab]**
  
- 2d convolution 
$$h[m, n] = \sum_{k, l} f[k, l] I[m - k, n - l]$$
  
`h=conv2(f, I) / h=imfilter(I, f, 'conv')` **[matlab]**

`conv2(I, f)` is the same as `filter2(rot90(f, 2), I)`

Correlation and convolution are identical when the filter is symmetric.

# Correlation and Convolution

- 2d correlation

$$h[m, n] = \sum_{k, l} f[k, l] I[m + k, n + l]$$

`h = cv2.filter2D(f, I)` [OpenCV]

- 2d convolution

$$h[m, n] = \sum_{k, l} f[k, l] I[m - k, n - l]$$

`h = cv2.filter2D(cv2.flip(f), I)` [OpenCV]

# Convolution properties

- Commutative:  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
  - But particular filtering implementations might break this equality, e.g., image edges
- Associative:  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
  - *Correlation* is **not** associative (rotation effect)
  - Why important?

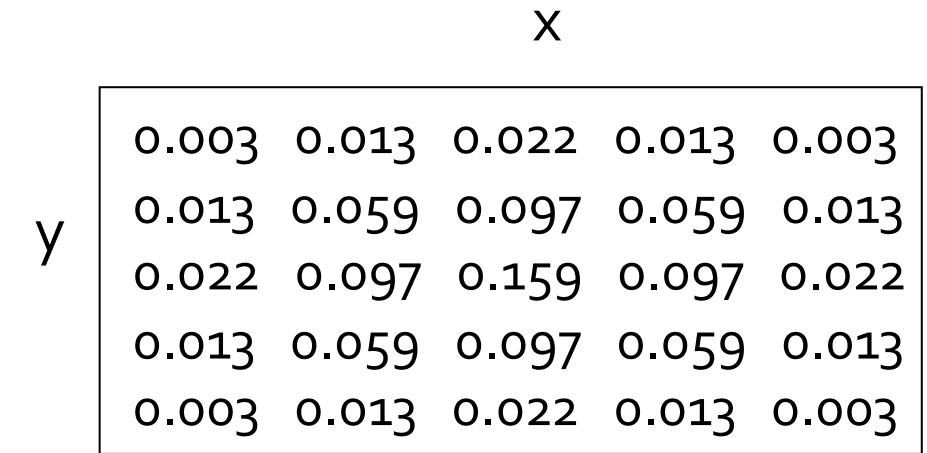
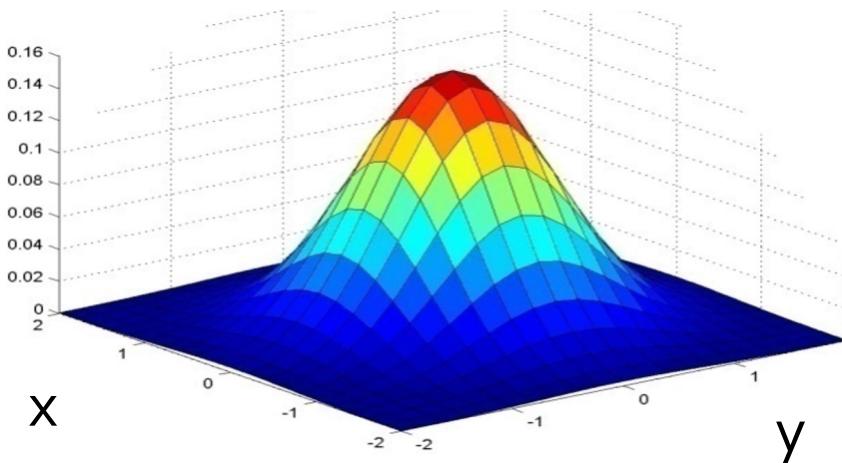
The mathematical symbol for convolution is the asterisk operator (\*) – not to be confused with multiplication in source code (\*)!

# Convolution properties

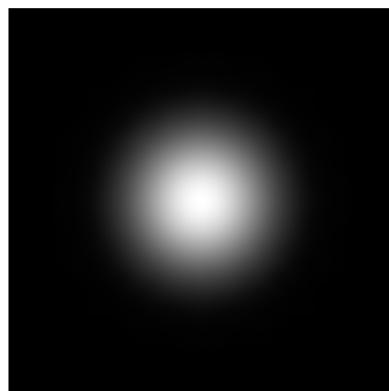
- Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out:  $ka * b = a * kb = k(a * b)$
- Identity: unit impulse  $e = [0, 0, 1, 0, 0] \Rightarrow a * e = a$

# Important filter: Gaussian

- Weight contributions of neighboring pixels by nearness

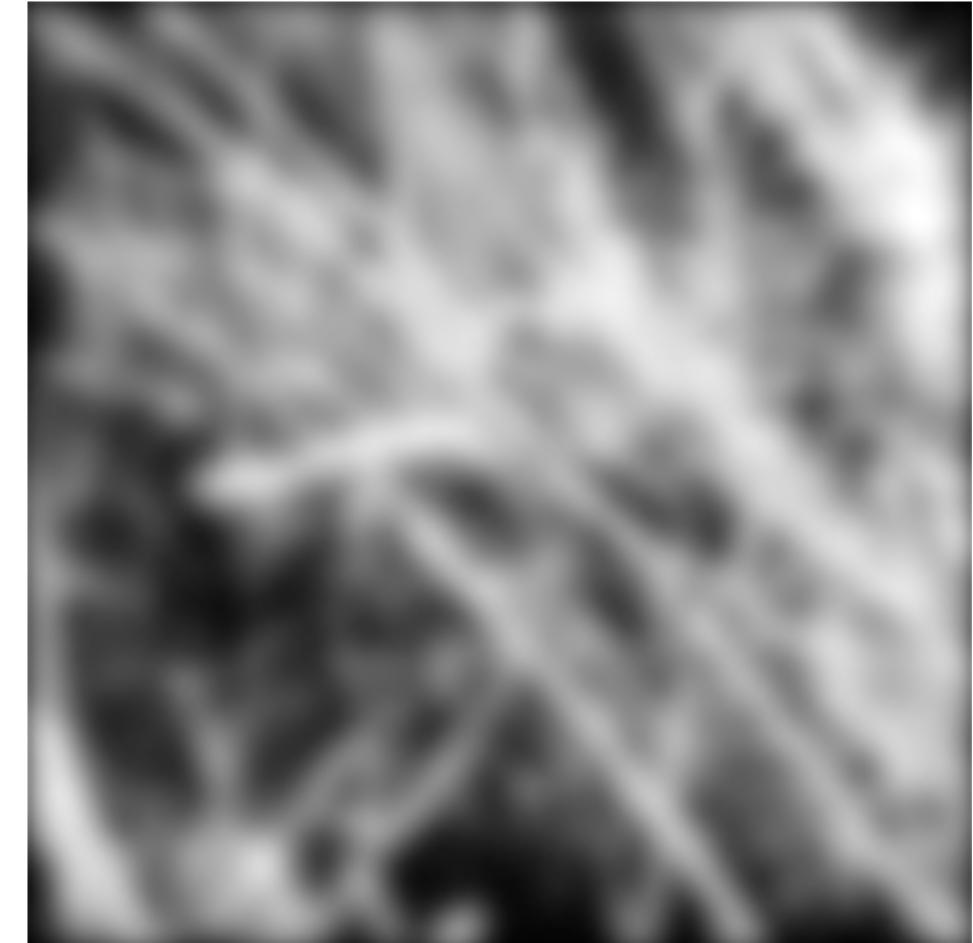


$5 \times 5, \sigma = 1$

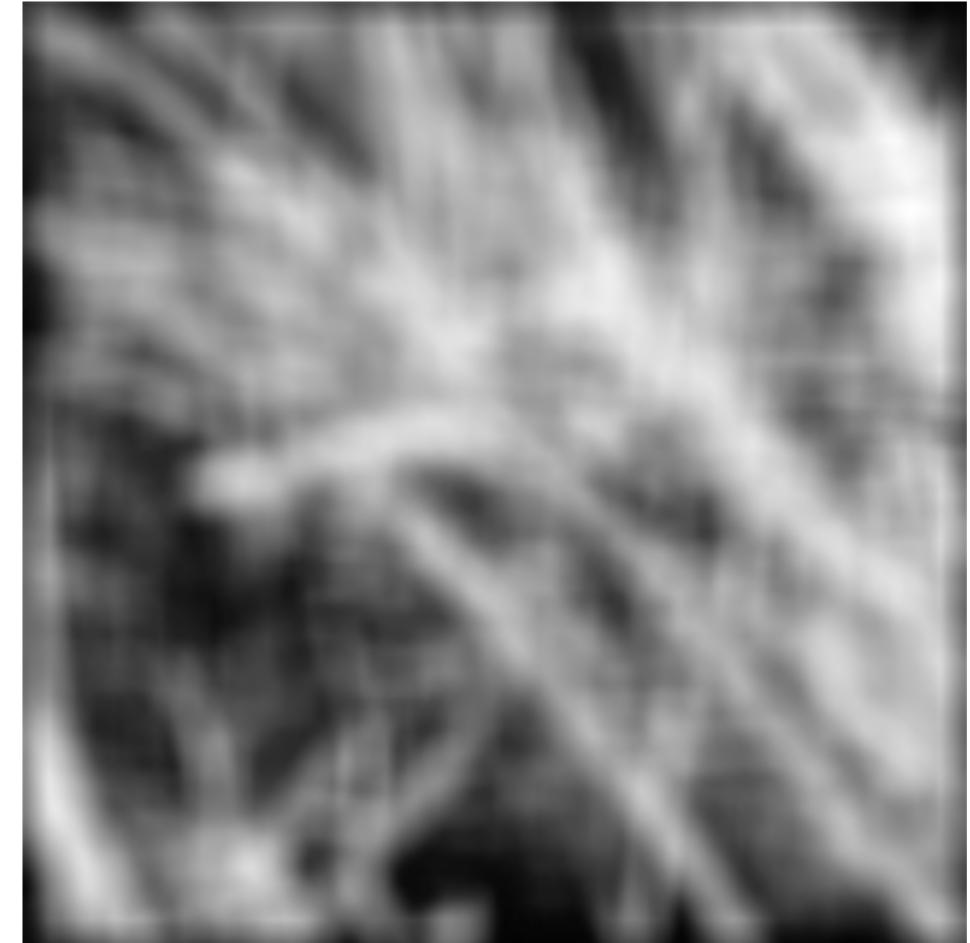


$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Smoothing with Gaussian filter



# Smoothing with box filter



# Properties of Gaussian filters

- Low-pass filter
  - Remove high-frequency components from the image
  - Images become more smooth
- Gaussian convolved with Gaussian...  
  ...is another Gaussian
- Convolving two times with Gaussian of width  $\sigma$  is same as convolving once with width  $\sigma\sqrt{2}$
- *Separable kernel*
  - Factors into product of two 1D Gaussians

## Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Separability: numerical example

2D convolution  
(center location only)

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & 65 & \\ \hline & & \\ \hline \end{array}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

# Separability: numerical example

Perform convolution  
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} & 11 & \\ & 18 & \\ & 18 & \end{matrix}$$

Followed by convolution  
along the remaining column:

$$\begin{matrix} 1 \\ 2 \\ 1 \end{matrix} * \begin{matrix} & 11 & \\ & 18 & \\ & 18 & \end{matrix} = \begin{matrix} & & \\ & & \\ & 65 & \end{matrix}$$

# Separability

Why is separability useful in practice?

MxN image, PxQ filter

- 2D convolution:  $\sim MNPQ$  multiply-adds (**computational complexity**)
- Separable 2D:  $\sim MN(P+Q)$  multiply-adds (**computational complexity**)

Speed up =  $PQ/(P+Q)$

9×9 filter =  $\sim 4.5 \times$  faster

# Practical matters

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



# Linear Filters – Examples



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

# Linear Filters – Examples



Original

0	0	0
0	0	1
0	0	0

?

# Linear Filters – Examples



Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

## Example: Box Filter

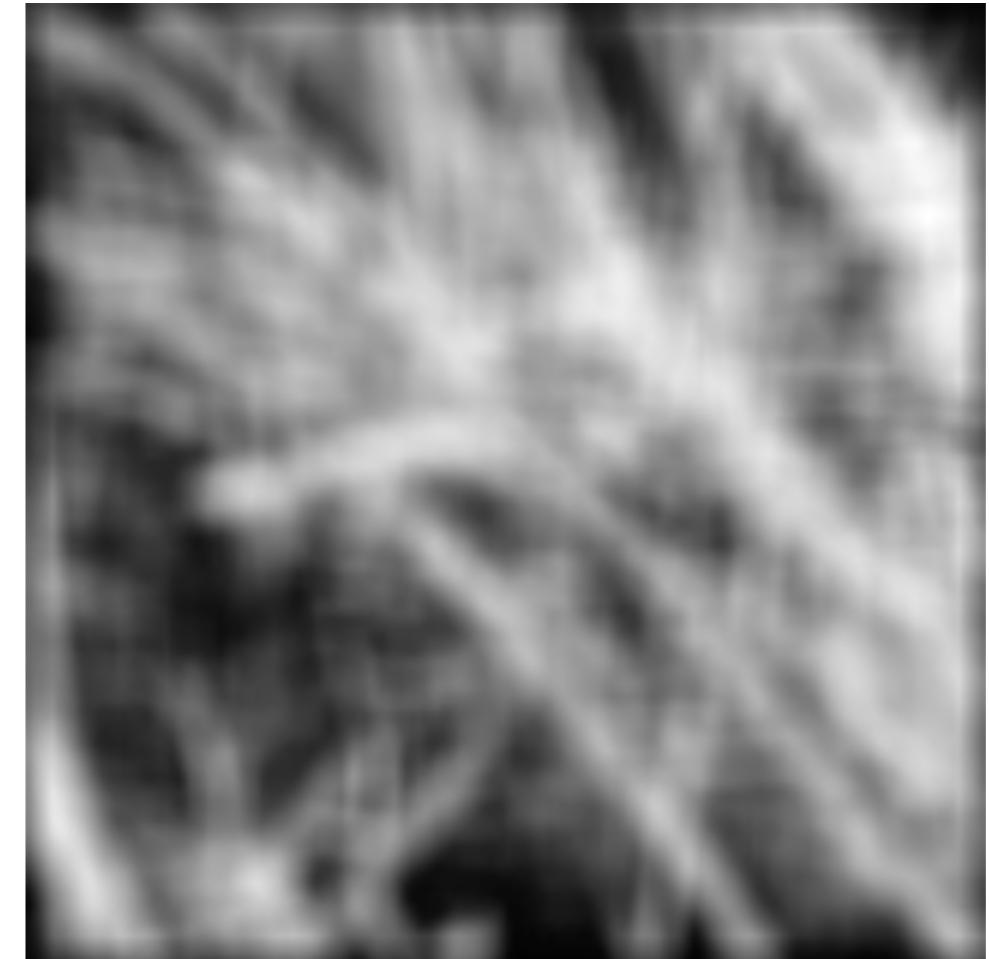
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)
- Why does it sum to one?

$f[\cdot, \cdot]$

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

# Smoothing with box filter

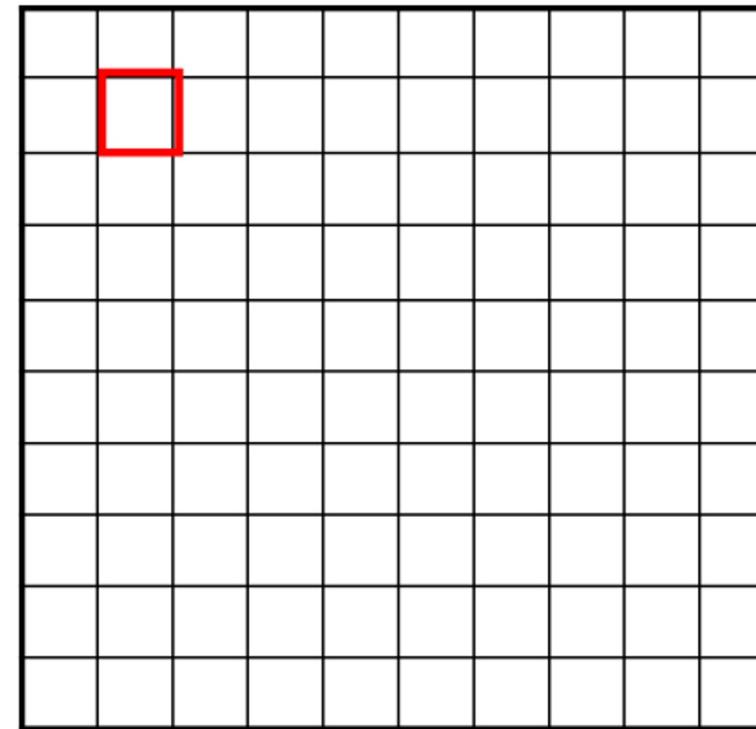


# Box Filter – Steps

$I[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$



$f[.,.]$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# Box Filter – Steps

$I[.,.]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10										

$f[.,.]$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# Box Filter – Steps

$I[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

	0	10	20	30	30				

$f[.,.]$

$\frac{1}{9}$			
1	1	1	
1	1	1	

# Box Filter – Steps

$I[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$h[.,.]$

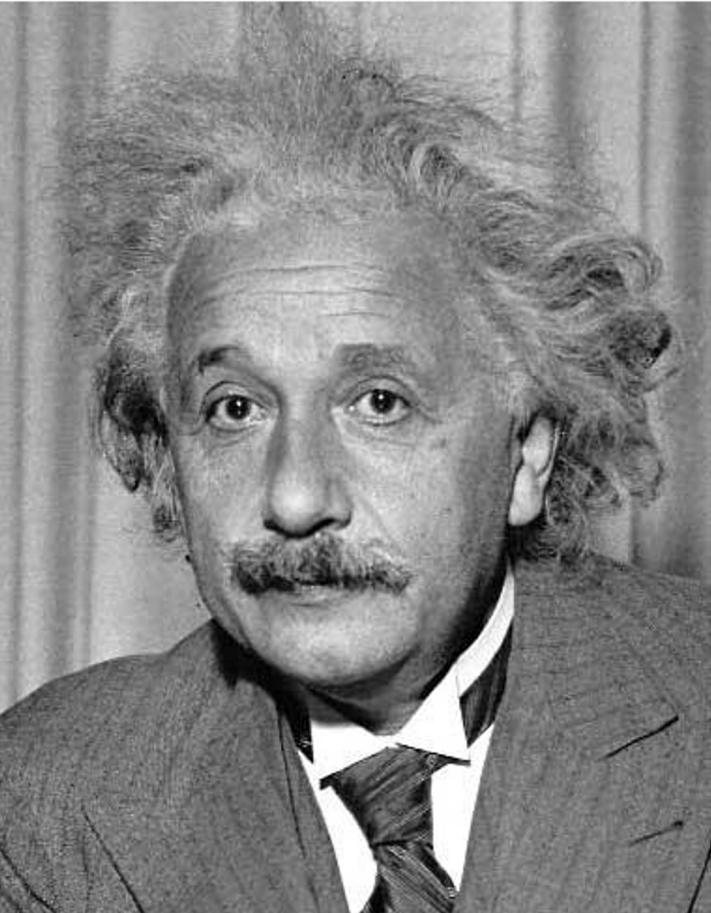
0	10	20	30	30	30	20	10	
0	20	40	60	60	60	40	20	
0	30	60	90	90	90	60	30	
0	30	50	80	80	90	60	30	
0	30	50	80	80	90	60	30	
0	20	30	50	50	60	40	20	
10	20	30	30	30	30	20	10	
10	10	10	0	0	0	0	0	

$f[.,.]$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

# Sobel filter



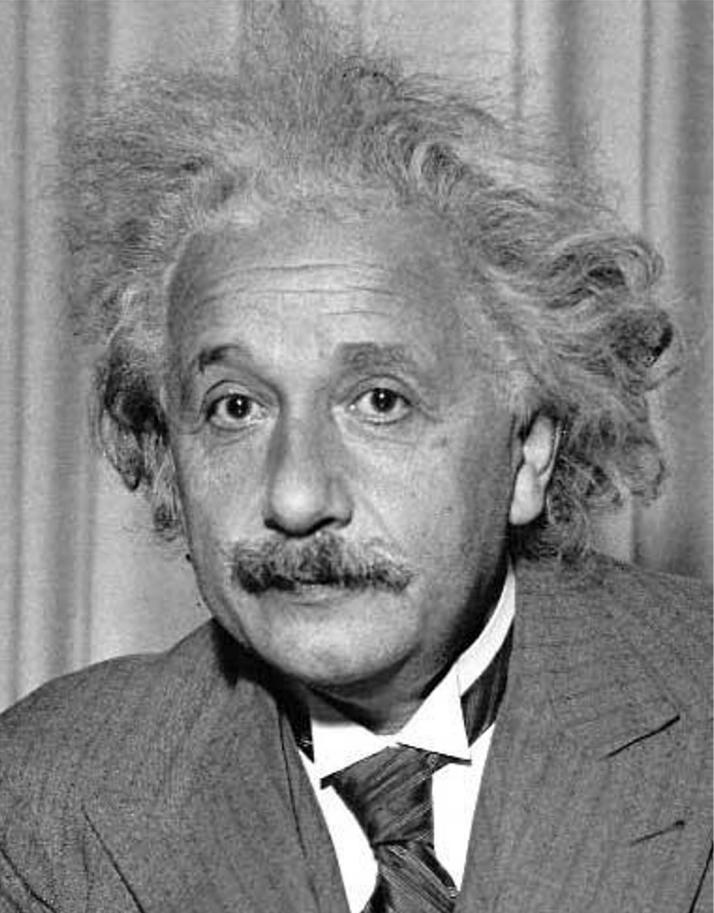
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge  
(absolute value)

# Sobel filter



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)

# Sobel filter

- Central differences
  - Edge enhancing
  - Slightly smoothing
  - Positive and negative output values

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

```
>> I = im2double( imread('luke.jpg') );  
>> h = imfilter(I, fspecial('sobel'));
```

# Sobel filter

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$



`imshow( h );`



`imshow( h + 0.5 );`



# Sobel filter

$h(:,:,1) < 0$



$h(:,:,1) > 0$

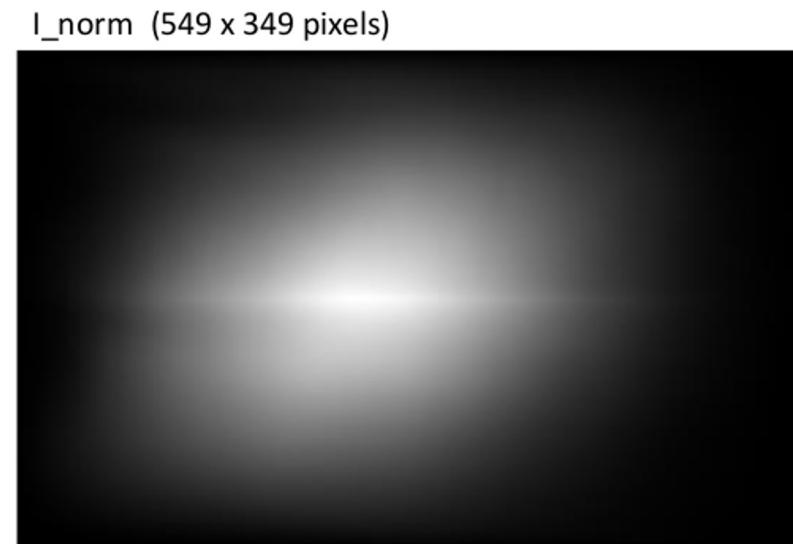
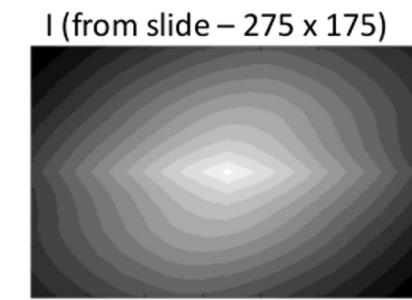
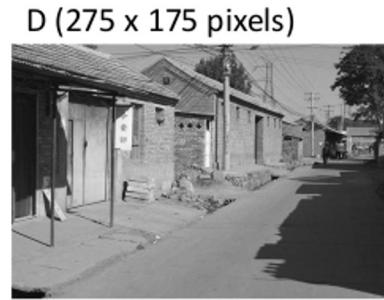
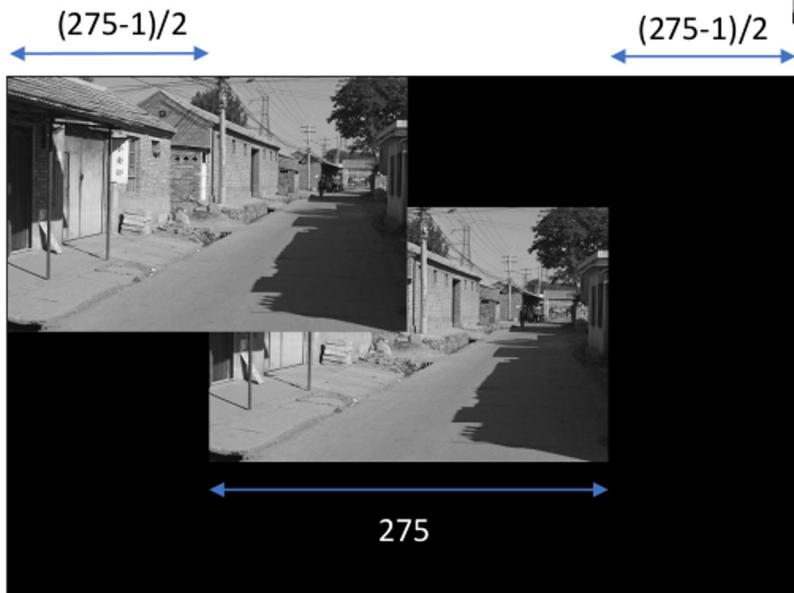


# Convolution in Convolutional Neural Networks

- Convolution is the basic operation in CNNs
- Learning convolution kernels allows us to learn which ‘features’ provide useful information in images.

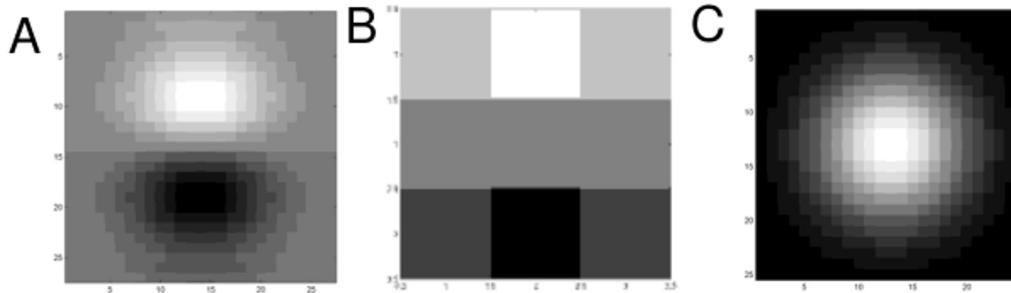
# What does it '*really*' mean to convolve?

$$I = D * D$$



# Convolution Vs. Correlation

$A = B * C$  – “because it kind of looks like it.”



C is a Gaussian filter  
(or something close to it),  
and we know that it ‘blurs’.

When the filter ‘looks like’ the image = ‘template matching’

Filtering viewed as comparing an image of  
what you want to find against all image regions.

# Convolution Vs. Correlation

subtract the mean

```
>> f = D[ 57:117, 107:167 ]  
>> f2 = f - np.mean(f)  
>> D2 = D - np.mean(D)
```

*Now zero centered.  
Score is higher only when dark parts  
match and when light parts match.*

```
>> I2 = correlate2d( D2, f2, 'same' )
```



# Convolution Vs. Correlation

What happens with convolution?

```
>> f = D[ 57:117, 107:167 ]
```

```
>> f2 = f - np.mean(f)
```

```
>> D2 = D - np.mean(D)
```

```
>> I2 = convolve2d( D2, f2, 'same' )
```

D2 (275 x 175 pixels)



f2  
61 x 61

I2

