

# Computer Vision

## (Summer Semester 2020)

Lecture 4, Part 1

Edge Detection

# Edges

- What types of edges do exist?
  - Edges in noisy images
  - Canny edge detector
- 
- Note: The core of these slides stems from the class CSCI 1430:  
“Introduction to Computer Vision” by James Tompkin, Fall 2017, Brown University.

# Low Level vs. High Level

Sensing

Data Representation

Edges

Corners

Descriptors

Camera Calibration

Alignment of Multi-view stereo

3D Reconstruction


Low



High

# Vanishing Point

- Any 2 parallel lines (real world/ground plane) have same **vanishing point** in the image plane
- Ray  $OV$ , is parallel to the parallel lines
- There can be multiple vanishing points




# Vanishing Point

- Any 2 parallel lines (real world/ground plane) have same **vanishing point** in the image plane
- Ray OV, is parallel to the parallel lines
- There can be multiple vanishing points




# Vanishing Lines

- Any set of parallel lines (ground plane) meet at a vanishing point
- Set of all the vanishing points form the horizon line OR the vanishing line
- Different planes define different vanishing lines




# Edge detection (Szeliski 4.2)

- **Goal:** Identify visual changes (discontinuities) in an image.
- **Why we care about edges?**
  - Recover viewpoint and geometry
  - Higher level vision tasks, e.g. for recognition




# Origin of Edges

**Edges are caused by a variety of factors**






# Closeup of edges




# Characterizing edges

An edge is a place of rapid change in the image intensity function


image



intensity function  
(along horizontal scanline)





first derivative





edges correspond to  
extrema of derivative

# Intensity profile





## With a little Gaussian noise



# Effects of noise

- Consider a single row or column of the image
  - Plotting intensity as a function of position gives a signal



**Where is the edge?**

# Effects of noise

- Difference filters respond strongly to noise
  - Image noise results in pixels that look very different from their neighbors
  - Generally, the larger the noise the stronger the response
- What can we do about it?

# Solution: smooth first


To find edges, look for peaks

in

$$\frac{d}{dx}(f * g)$$

$$f * g$$

$$\frac{d}{dx}(f * g)$$



# Derivative theorem of convolution


- Convolution is differentiable:  $\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$

- This saves us one operation:

$f$


$\frac{d}{dx}g$

$f * \frac{d}{dx}g$







# Derivative of 2D Gaussian filter




$$* \begin{bmatrix} 1 & -1 \end{bmatrix} =$$



# Tradeoff between smoothing and localization



1 pixel



3 pixels





7 pixels

- Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”.

# Designing an edge detector

- Criteria for a good edge detector:
  - **Good detection:** the optimal detector should find all real edges, ignoring noise or other artifacts
  - **Good localization**
    - the edges detected must be as close as possible to the true edges
    - the detector must return one point only for each true edge point
- Cues of edge detection
  - Differences in color, intensity, or texture across the boundary
  - Continuity and closure
  - High-level knowledge

# Closeup of edges



# Designing an edge detector




- “All real edges”
  - We can aim to differentiate later on which edges are ‘useful’ for our applications.
  - If we can’t find all things which *could* be called an edge, we don’t have that choice.
- Is this possible?

# Where do humans see boundaries?

human segmentation

gradient magnitude


image



- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

# 45 years of boundary detection



*Arbelaez, Maire, Fowlkes, and Malik.  
TPAMI 2011*

# State of edge detection

- Local edge detection works well
  - ‘False positives’ from illumination and texture edges (depends on our application).
- Some methods to take into account longer contours
- Modern methods that actually “learn” from data. *[not in this class]*
- Poor use of object and high-level information.



# Canny edge detector

- Probably the most widely used edge detector in computer vision.
- Theoretical model: step-edges corrupted by additive Gaussian noise.
- Canny showed that first derivative of Gaussian closely approximates the operator that optimizes the product of *signal-to-noise ratio* and localization.

*J. Canny, **A Computational Approach To Edge Detection**, IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.*

# Demonstrator Image


`rgb2gray('img.png')`





# Canny edge detector

1. Filter image with x, y derivatives of Gaussian


# Derivative of Gaussian filter



x-direction



y-direction




# Compute Gradients

X Derivative of Gaussian



Y Derivative of Gaussian



(x2 + 0.5 for visualization)





# Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient

# Compute Gradient Magnitude



$\text{sqrt}( \text{XDerivOfGaussian}.^2 + \text{YDerivOfGaussian}.^2 ) =$  gradient magnitude




(x4 for visualization)

# Compute Gradient Orientation

- Threshold magnitude at minimum level
- Get orientation via  **$\theta = \text{atan2}(g_y, g_x)$**

**Thresholded:**






# Compute Gradient Orientation

- Threshold magnitude at minimum level
- Get orientation via  $\text{theta} = \text{atan2}(gy, gx)$


Unthresholded:



# Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
  - Thin multi-pixel wide “ridges” to single pixel width


# Non-maximum suppression for each orientation



At pixel q:

We have a maximum if the value is larger than those at both p and at r.

Interpolate along gradient direction to get these values.



# Before Non-max Suppression



Gradient magnitude (x4 for visualization)

# After non-max suppression




Gradient magnitude (x4 for visualization)

# Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
  - Thin multi-pixel wide “ridges” to single pixel width
4. ‘Hysteresis’ Thresholding



## ‘Hysteresis’ thresholding

- Two thresholds – high and low
- Grad. mag.  $>$  high threshold? = strong edge
- Grad. mag.  $<$  low threshold? noise
- In between = weak edge
- ‘Follow’ edges starting from strong edge pixels
- Continue them into weak edges
- Connected components (Szeliski 3.3.4)




# Final Canny Edges

$$\sigma = \sqrt{2}, t_{low} = 0.05, t_{high} = 0.1$$






## Effect of $\sigma$ (Gaussian kernel spread/size)



Original



$\sigma = \sqrt{2}$



$\sigma = 4\sqrt{2}$

The choice of  $\sigma$  depends on desired behavior

- large  $\sigma$  detects large scale edges
- small  $\sigma$  detects fine features

# Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
  - Thin multi-pixel wide “ridges” to single pixel width
4. ‘Hysteresis’ Thresholding:
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them
  - ‘Follow’ edges starting from strong edge pixels
  - Connected components (Szeliski 3.3.4)

**MATLAB:** `edge(image, ‘canny’)`