

Computer Vision

(Summer Semester 2020)

Lecture 5, Part 2

Feature Descriptors (2)

Feature Descriptors and Matching

- SIFT: scale-invariant image descriptor
- Feature matching

Note: The core of these slides stems from the class CSCI 1430: “Introduction to Computer Vision” by James Tompkin, Fall 2017, Brown University.

Review: Basic Feature Detection Algorithm

1. Compute horizontal and vertical derivatives, viz. I_x and I_y

$$I_x = I * \text{LoG}_x \quad \& \quad I_y = I * \text{LoG}_y$$

2. Compute the moment matrix \mathbf{M}
$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

3. Convolve each of these images with a larger gaussian

4. Compute thresholds to find the featureness (eg: cornerness) scores

5. Find local maxima above a threshold and report them as features (eg:

MSER)



Local Image Descriptors (Szeliski 4.1)

Links


- <https://www.robots.ox.ac.uk/~vgg/research/affine/>
- [Distinctive Image Features from Scale-Invariant Keypoints](#)
- <https://aishack.in/tutorials/sift-scale-invariant-feature-transform-introduction/>

Local features: main components

- 1) Detection:
Find a set of distinctive key points.
- 2) Description:
Extract feature descriptor around each interest point as vector.



$$\mathbf{x}_1 = [x_1^{(1)}, \square, x_d^{(1)}]$$



$$\mathbf{x}_2 = [x_1^{(2)}, \square, x_d^{(2)}]$$

- 3) Matching:
Compute distance between feature vectors to find correspondence.

$$d(\mathbf{x}_1, \mathbf{x}_2) < T$$






Image Representations: Histograms



Global histogram to represent distribution of features

- Color, texture, depth, oriented gradients...

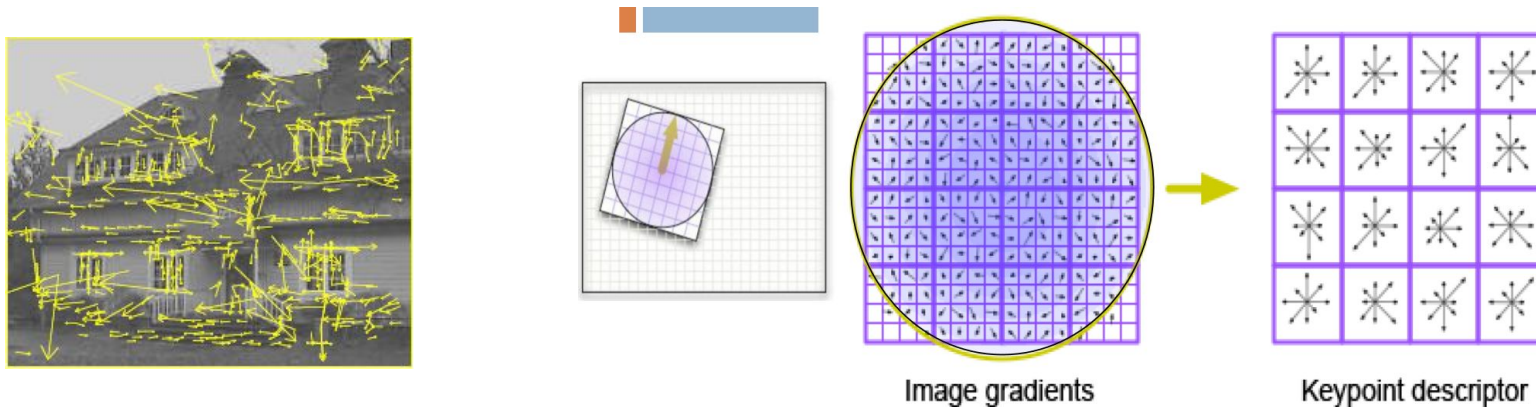
Local histogram per detected point

SIFT – Lowe IJCV 2004

- Compute local histograms of oriented gradients
- Find Difference of Gaussian scale-space extrema (for target scale)
- Post-processing
 - Position interpolation
 - Discard low-contrast points
 - Eliminate points along edges
- Orientation estimation
- Descriptor extraction
 - Motivation: We want some sensitivity to spatial layout, but not too much, so blocks of histograms give us that.



SIFT Descriptor Extraction

- Given a keypoint with scale and orientation:
- Pick scale-space image which most closely matches estimated scale
- Resample image to match orientation OR
- Subtract detector orientation from vector to give invariance to general image rotation.




SIFT Orientation Normalization

- Compute orientation histogram
- Select dominant orientation Θ
- Normalize: rotate to fixed orientation



SIFT Descriptor Extraction


- Given a keypoint with scale and orientation



SIFT Descriptor Extraction


- Within each 4x4 window

Gradient
magnitude
and
orientation



8 bin 'histogram'
- add magnitude
amounts!

Weight magnitude
that is added to
'histogram' by
Gaussian




SIFT Descriptor Extraction

- Extract 8 x 16 values into 128-dim vector
- Illumination invariance:
 - Working in gradient space, so robust to $I = I + b$
 - Normalize vector to $[0...1]$
 - Robust to $I = \alpha I$ brightness changes
 - Clamp all vector values > 0.2 to 0.2.
 - Robust to “non-linear illumination effects”
 - Image value saturation / specular highlights
 - Renormalize

Implementation Notes

- Efficient Implementation
 - Filter using oriented kernels based on directions of histogram bins.
 - Called 'steerable filters'
- Sources
 - tutorial:
<http://aishack.in/tutorials/sift-scale-invariant-feature-transform-features/>
 - Lowe's original paper: <http://www.cs.ubc.ca/~lowe/papers/ijcv04.pdf>

Local Descriptors: SURF




- Fast approximation of SIFT idea
 - Efficient computation by 2D box filters & integral images
⇒ 6 times faster than SIFT
 - Equivalent quality for object identification
- GPU implementation available
 - Feature extraction @ 200Hz
(detector + descriptor, 640×480 img)
 - <http://www.vision.ee.ethz.ch/~surf>

[Bay, ECCV'06], [Cornelis, CVGPU'08]

Review: Local Descriptors

- Most features can be thought of as templates, histograms (counts), or combinations
- The ideal descriptor should be
 - Robust and Distinctive
 - Compact and Efficient
- Most available descriptors focus on edge/gradient information
 - Capture texture information
 - Color rarely used





Available at a web site near you...

- Many local feature detectors have executables available online:
 - <http://www.robots.ox.ac.uk/~vgg/research/affine>
 - <http://www.cs.ubc.ca/~lowe/keypoints/>
 - <http://www.vision.ee.ethz.ch/~surf>


Feature Matching (Szeliski 4.1)

Local features: main components

- 1) Detection:
Find a set of distinctive key points.
- 2) Description:
Extract feature descriptor around each interest point as vector.



$$\mathbf{x}_1 = [x_1^{(1)}, \square, x_d^{(1)}]$$




$$\mathbf{x}_2 = [x_1^{(2)}, \square, x_d^{(2)}]$$

- 3) **Matching:**
Compute distance between feature vectors to find correspondence.


$$d(\mathbf{x}_1, \mathbf{x}_2) < T$$



How do we decide which features match?



Distance: 0.34, 0.30, 0.40



Distance: 0.61, 1.22

Feature Matching

- Criteria:
 - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
 - Match point to lowest distance (nearest neighbor)
 - Ignore anything higher than threshold (no match!)
- Problems:
 - Threshold is hard to pick
 - Non-distinctive features could have lots of close matches, only one of which is correct

Nearest Neighbor Distance Ratio

Compare distance of closest (NN1) and second-closest (NN2) feature vector neighbor.

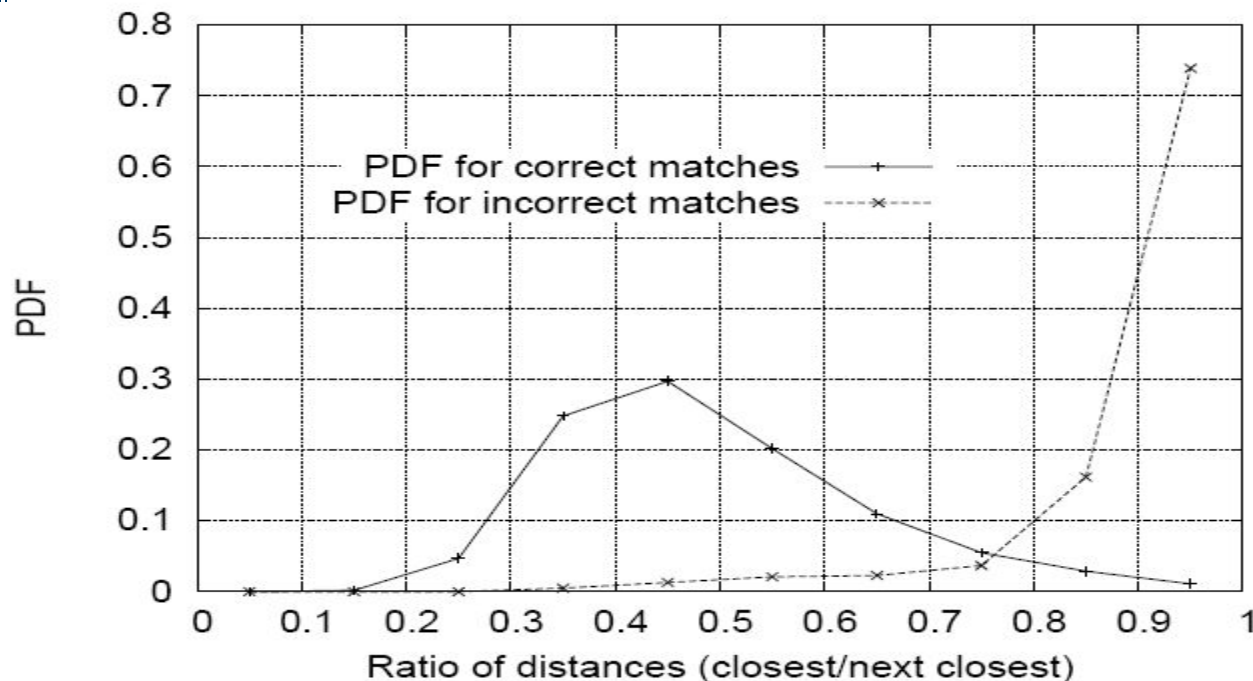
- If $NN1 \approx NN2$, ratio $\frac{NN1}{NN2}$ will be ≈ 1 -> matches too close.
- As $NN1 \ll NN2$, ratio $\frac{NN1}{NN2}$ tends to 0.

Sorting by this ratio puts matches in order of confidence.

Threshold ratio – but how to choose?

Nearest Neighbor Distance Ratio

- Lowe computed a probability distribution functions of ratios
- 40,000 keypoints with hand-labeled ground truth



Ratio threshold depends on your application's view on the trade-off between the number of false positives and true positives!