Friedrich-Alexander-Universität
Technische Fakultät

Machine Learning
Data Analytics

FAU

Machine Learning for Time Series

(MLTS or MLTS-Deluxe Lectures)

Dr. Dario Zanca

Machine Learning and Data Analytics (MaD) Lab
Friedrich-Alexander-Universität Erlangen-Nürnberg
25.10.2022

Organisational Information

Lectures (online)

A new lecture recording every Tuesday.
Consultation hours from November 15th, h. 8:15 – 9:30

Exercises (online)

Live Zoom Session starting on November 9th

Recordings will be uploaded

Project (online)

Introduction during the first exercise Live Zoom Session (November 9th)

Applications started on Nov 9th

Topics overview

- Time series fundamentals and definitions (2 lectures)
- Bayesian Inference (1 lecture) ←
- Gussian processes (2 lectures)
- State space models (2 lectures)
- Autoregressive models (1 lecture)
- Data mining on time series (1 lecture)
- Deep learning on time series (4 lectures)
- Domain adaptation (1 lecture)

In this lecture…

1. Bayes Theorem
2. Bayesian Model Selection
3. Prior Distributions
4. Linear Regression (Bayesian treatment)

Bayesian Inference
Bayes' Theorem

Bayes' Theorem

Formulation

The **Bayes' Theorem** was formulated by the English philosopher **Thomas Bayes** (1701 – 1761), whose notes were edited and published posthumously by Richard Price.

Bayes' theorem is stated mathematically as the following equation:

$$\mathrm{P}(A|B) = \frac{\mathrm{P}(B|A)\mathrm{P}(A)}{\mathrm{P}(B)}$$

where A and B are events, and $P(B) \neq 0$.

Portrait from: Terence O'Donnell, *History of Life Insurance in Its Formative Years* (Chicago: American Conservation Co:, 1936), p. 335

Bayes' Theorem

Formulation

Likelihood
The probability of event B occurring given that A is true

Posterior probability
The probability of event A occurring given that B is true

$$P(A|B) = \frac{P(B|A)\,P(A)}{P(B)}$$

Marginal probability
The probability of observing B without any given conditions

Prior probability
The probability of observing A without any given conditions

Portrait from: Terence O'Donnell, *History of Life Insurance in Its Formative Years* (Chicago: American Conservation Co:, 1936), p. 335

Bayes' Theorem

An example. Iterative application of the Bayes' theorem.

We know that:
- Disease chance: 1%
- Test accuracy: 95%

$\boldsymbol{A}$: Having the disease
$\boldsymbol{B}$: Testing positive to the disease

$\mathrm{P}(\boldsymbol{B}|\boldsymbol{A})$: Test accuracy (Likelihood)
$\mathrm{P}(\boldsymbol{B})$: Prob. of a positive test (Marginal)
$\mathrm{P}(\boldsymbol{A})$: Disease chance (Prior)

First positive test

$$\mathrm{P}(A|B) = \frac{.95\times 0.01}{0.95 \cdot 0.01 + (1-0.95)(1-0.01)} = 0.161$$

Second positive test

$$\mathrm{P}(A|B) = \frac{.95\times 0.161}{0.95 \cdot 0.161 + (1-0.95)(1-0.161)} = 0.785$$

Third positive test

$$\mathrm{P}(A|B) = \frac{.95\times 0.785}{0.95 \cdot 0.785 + (1-0.95)(1-0.785)} = 0.986$$

$$\mathrm{P}(A|B) = \frac{\mathrm{P}(B|A)\,\mathrm{P}(A)}{\mathrm{P}(B)}$$

Portrait from: Terence O'Donnell, *History of Life Insurance in Its Formative Years* (Chicago: American Conservation Co:, 1936), p. 335

Bayesian modelling

Example of a linear model

Let $\mathcal{D}$ denote the **observed data**,

$$\mathcal{D} = \{x^{(n)}, y^{(n)}\}$$

with $x^{(n)} \in \mathcal{R}$ represents the input, and $y^{(n)} \in \mathcal{R}$ represents the output (labels).

The **model** is defined as

$$y^{(n)} = \omega_0 + \omega_1 x^{(n)} + \omega_2 x^{(n)} \ldots + \omega_2 x^{(n)} + \epsilon$$

where data noise is Gaussian distributed, i.e., $\epsilon \sim \mathcal{N}(0, \sigma^2)$.

We denote with θ the **unknown parameters**,

$$\theta = (\omega_0, \ldots, \omega_m, \sigma)$$

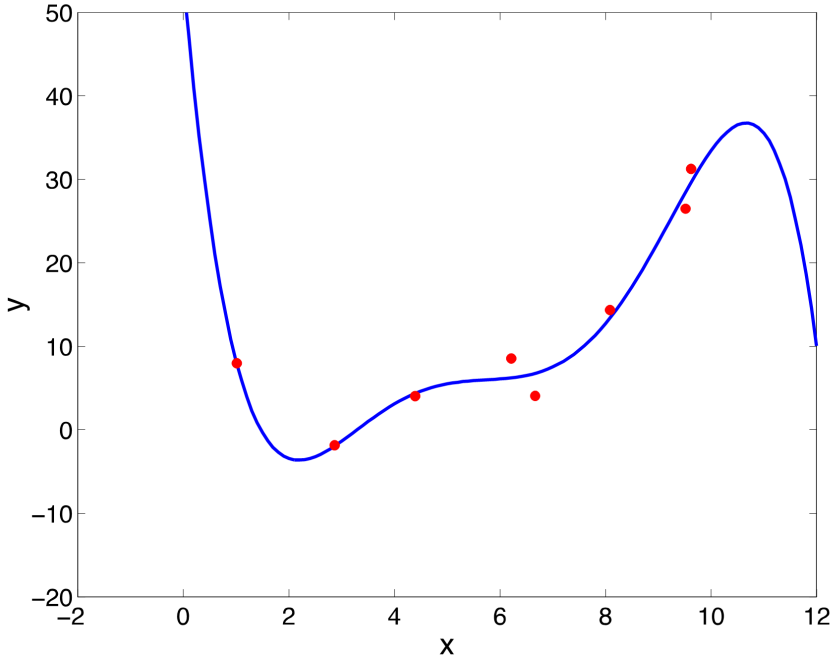

Bayesian modelling

Example of a linear model

Data: $\mathcal{D} = \{x^{(n)}, y^{(n)}\}$

Model: $y^{(n)} = \omega_0 + \omega_1 x^{(n)} + \omega_2 x^{(n)} \ldots + \omega_2 x^{(n)} + \epsilon$

Unknown parameters: $\theta = (\omega_0, \ldots, \omega_m, \sigma)$

Goal: To infer $\boldsymbol{\theta}$ from the data and to predict future outputs $p(y|x, \boldsymbol{\theta}, \boldsymbol{\mathcal{D}})$

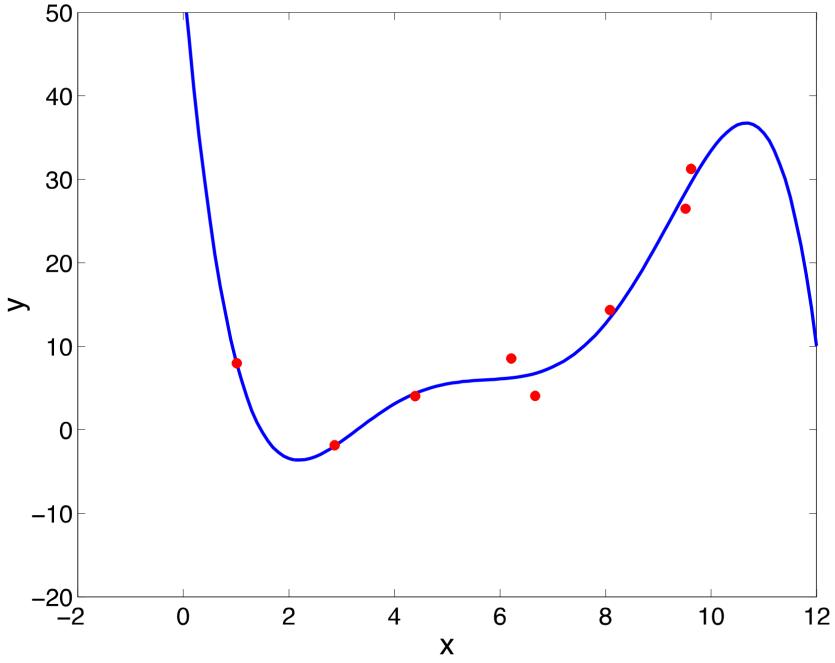

Bayesian modelling

$p(\boldsymbol{\mathcal{D}}|\boldsymbol{\theta})$: likelihood of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta})$: prior probability of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta}|\boldsymbol{\mathcal{D}})$: posterior of $\boldsymbol{\theta}$ given $\boldsymbol{\mathcal{D}}$

$p(\boldsymbol{\mathcal{D}})$: marginal probability of $\boldsymbol{\mathcal{D}}$

$p(y|x, \boldsymbol{\mathcal{D}})$: predictive distribution

Bayesian modelling

$p(\boldsymbol{\mathcal{D}}|\boldsymbol{\theta})$: likelihood of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta})$: prior probability of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta}|\boldsymbol{\mathcal{D}})$: posterior of $\boldsymbol{\theta}$, given $\boldsymbol{\mathcal{D}}$

$p(\boldsymbol{\mathcal{D}})$: marginal probability of $\boldsymbol{\mathcal{D}}$

$p(y|x, \boldsymbol{\mathcal{D}})$: predictive distribution

Bayes' rule:

$$p(\theta|\boldsymbol{\mathcal{D}}) = \frac{p(\boldsymbol{\mathcal{D}}|\theta)p(\theta)}{P(\boldsymbol{\mathcal{D}})}$$

Prediction of a new point:

$$p(y|x, \boldsymbol{\mathcal{D}}) = \int p(y|\theta, x, \boldsymbol{\mathcal{D}})p(\theta|\boldsymbol{\mathcal{D}})\, d\theta$$

Bayesian modelling

Observations

- In contrast to the maximum likelihood estimation (MLE), in Bayesian learning **we average over possible parameter settings** rather than optimizing over parameter space.
- Bayesian inference gives us a **systematic way to express our uncertainty** about future predictions. Prediction is not just a point estimate (as for MLE) but has a probability form that expresses the uncertainty about the predictions.

Bayesian Inference

Bayesian Model Selection

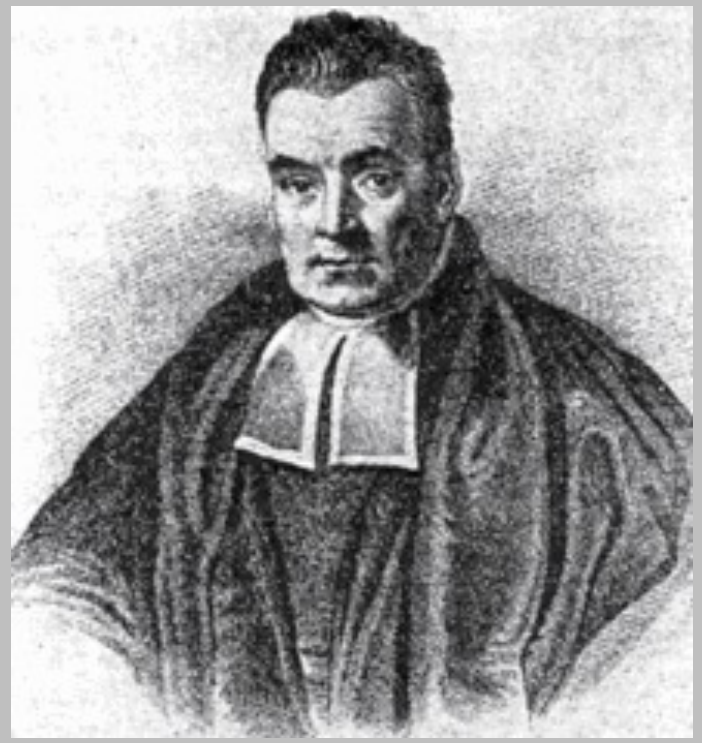


ould not be

The principle of Occam’s Razor

The **principle of Occam’s razor** in its original formulation states that:

"Entia non sunt multiplicanda praeter necessitatem"

(In English, “Entities should not be multiplied unnecessarily”)

Many scientists have adopted or reformulated the Occam's Razor principle, which is often cited in stronger forms, as in the following statement:

- "If you have two theories that both explain the observed facts, then you should use the simplest until more evidence comes along"
- “One should pick the simplest model that adequately explains the data”

Picture from: https://www.britannica.com/topic/Occams-razor

The model Selection problem

We could perform K-fold cross-validation (CV) to estimate the generalization error of all candidates model.

However, it requires fitting each candidate model K times!

→ A more efficient approach is given by Bayesian modelling

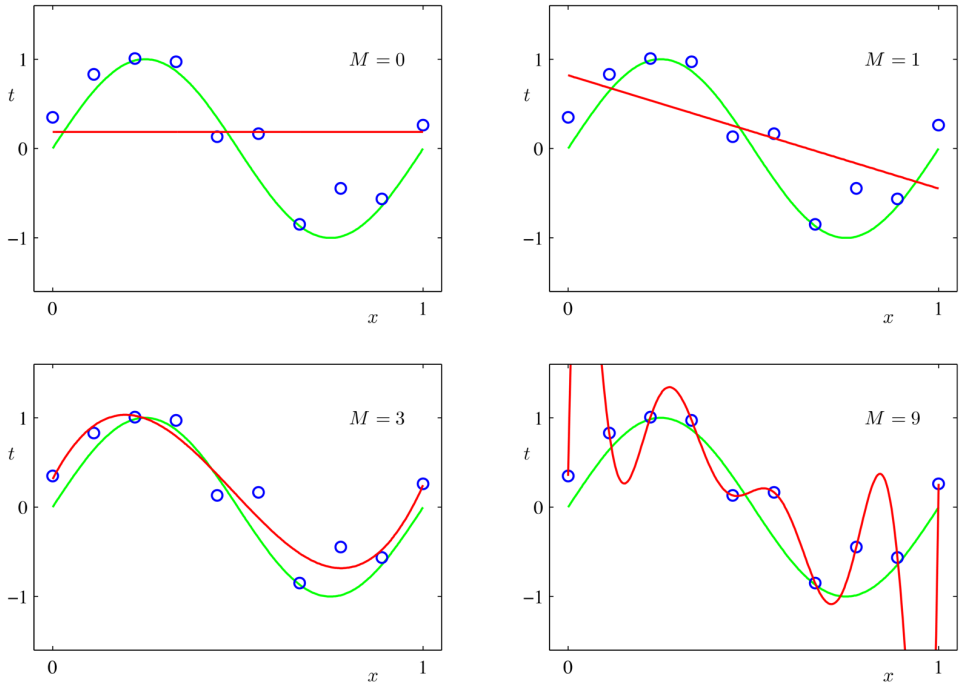

Which of the above models represents data the best?

Occam's razor and Bayesian model selection

We can compare different models using the marginal likelihood:

$$p(\mathcal{D}|M_i) = \int p(\mathcal{D}|\theta, M_i)p(\theta|M_i)\, d\theta$$

- Model classes that are **too simple** are unlikely to generate the data set.
- Model classes that are **too complex** can generate many possible data sets, so again, they are unlikely to generate that particular data set $\mathcal{D}$.

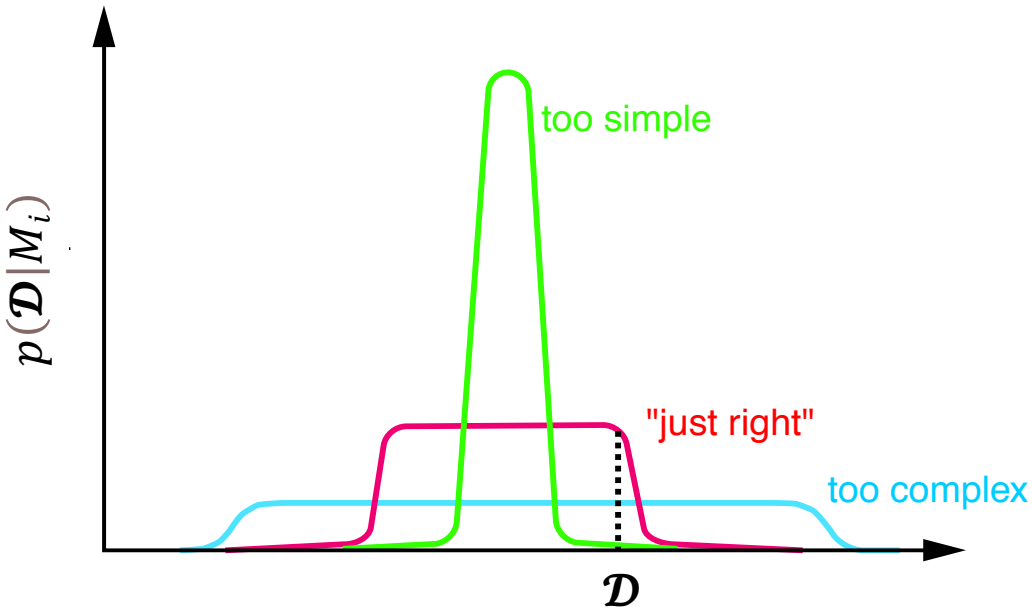

Image from: Rasmussen, C., & Ghahramani, Z. (2000). Occam's razor. Advances in neural information processing systems, 13.

Occam's razor and Bayesian model selection

To understand the Bayesian Occam's razor, we notice that:

$$\int_{\mathcal{D}} p(\mathcal{D}|M_i) = 1$$

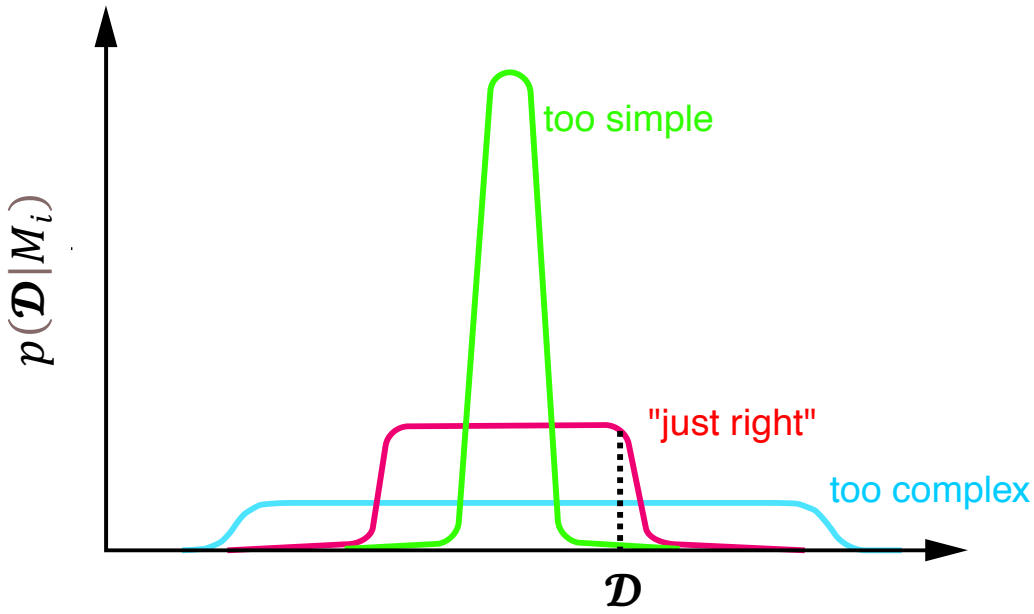

Image from: Rasmussen, C., & Ghahramani, Z. (2000). Occam's razor. Advances in neural information processing systems, 13.

Occam's razor and Bayesian model selection

To understand the Bayesian Occam's razor, we notice that:

$$\int_{\mathcal{D}} p(\mathcal{D}|M_i) = 1$$

Intuitively, complex models which can predict *many* datasets, must spread their probability mass
→ They don't attribute large probability for any given data set as simpler models.

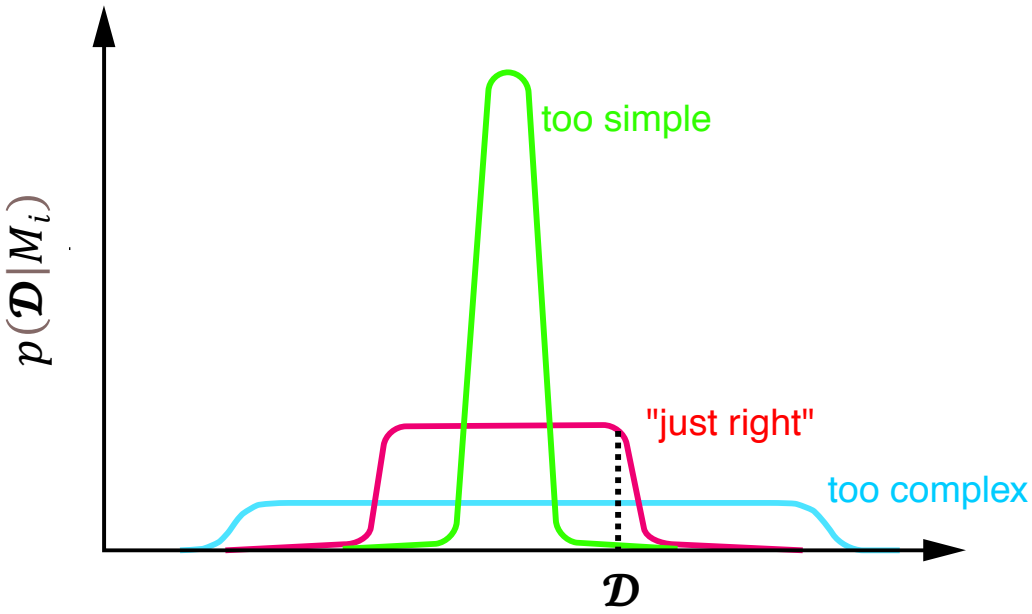

Image from: Rasmussen, C., & Ghahramani, Z. (2000). Occam's razor. Advances in neural information processing systems, 13.

Bayesian Occam's razor

A concrete example

We plot polynomials of degrees 1, 2 and 3 fit to N=5 data points using (empirical) Bayes.

True function

Prediction

$\pm\sigma$ around the mean

There is not enough data to justify a complex model, so the best model is d = 1.

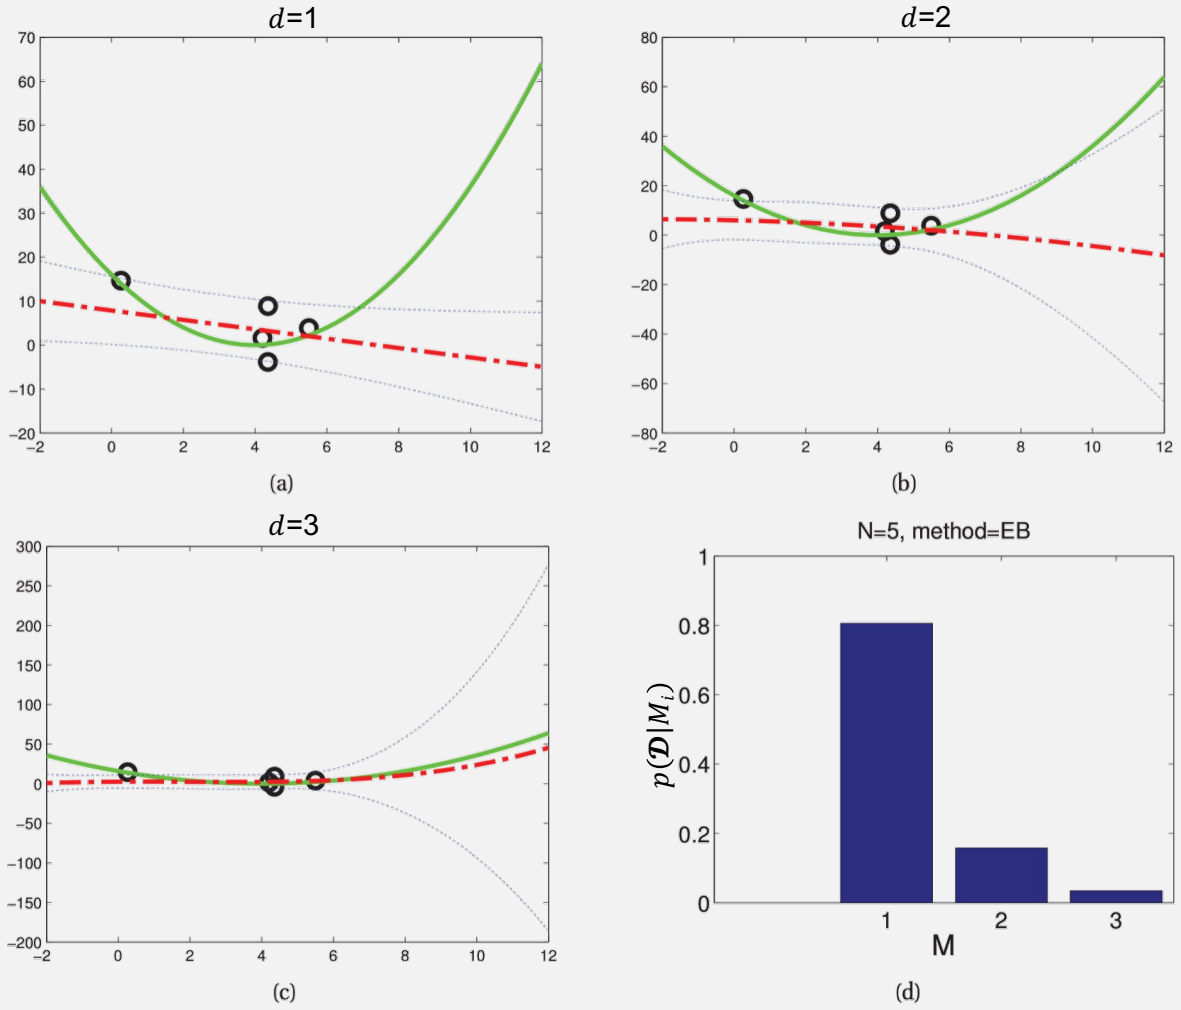

Bayesian Occam's razor

A concrete example

We plot polynomials of degrees 1, 2 and 3 fit to N=30 data points using (empirical) Bayes.

True function

Prediction

$\pm\sigma$ around the mean

When more data is available, d = 2 is the right model

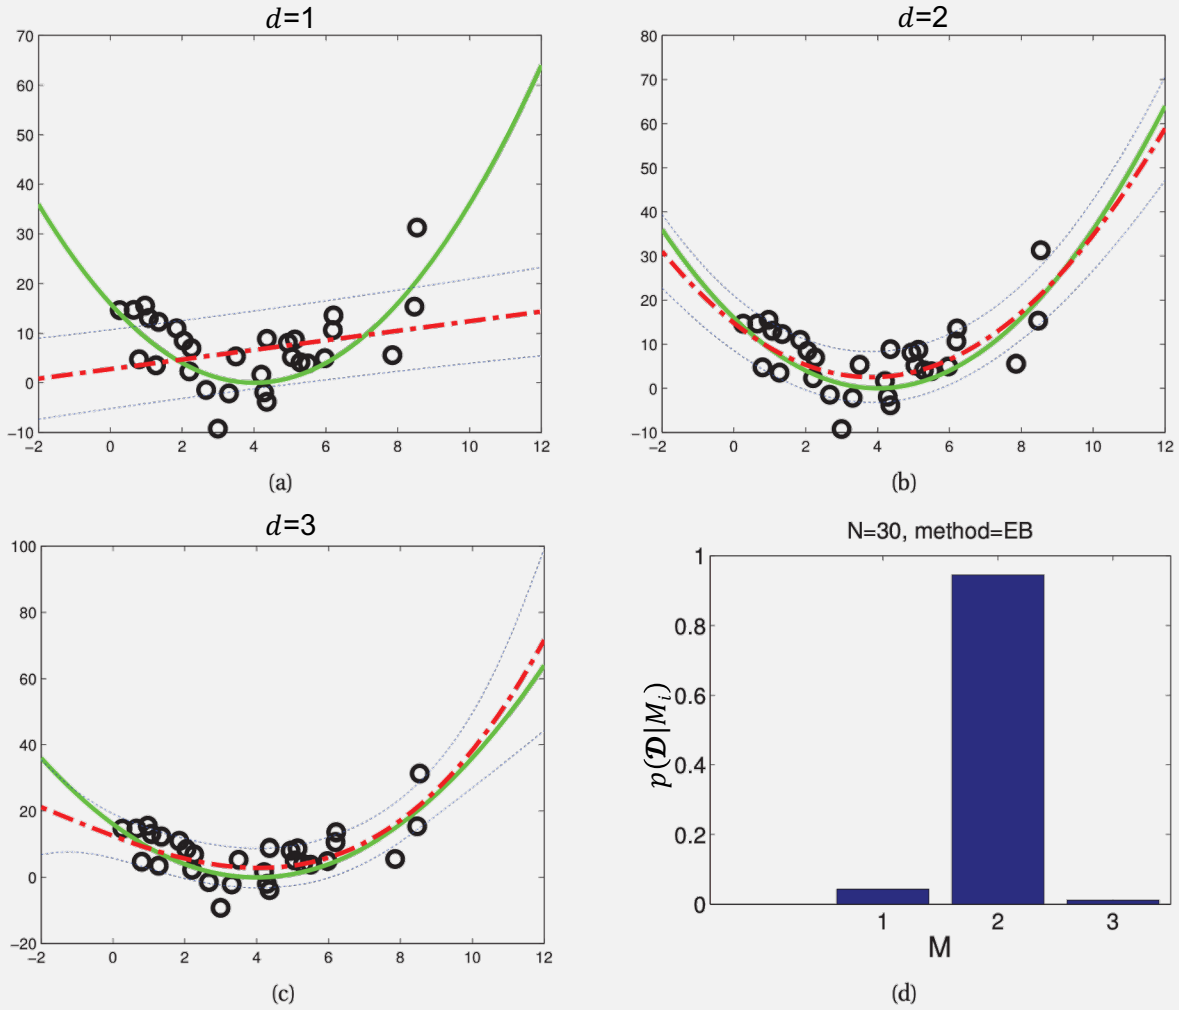

Bayesian Inference
Prior Distributions

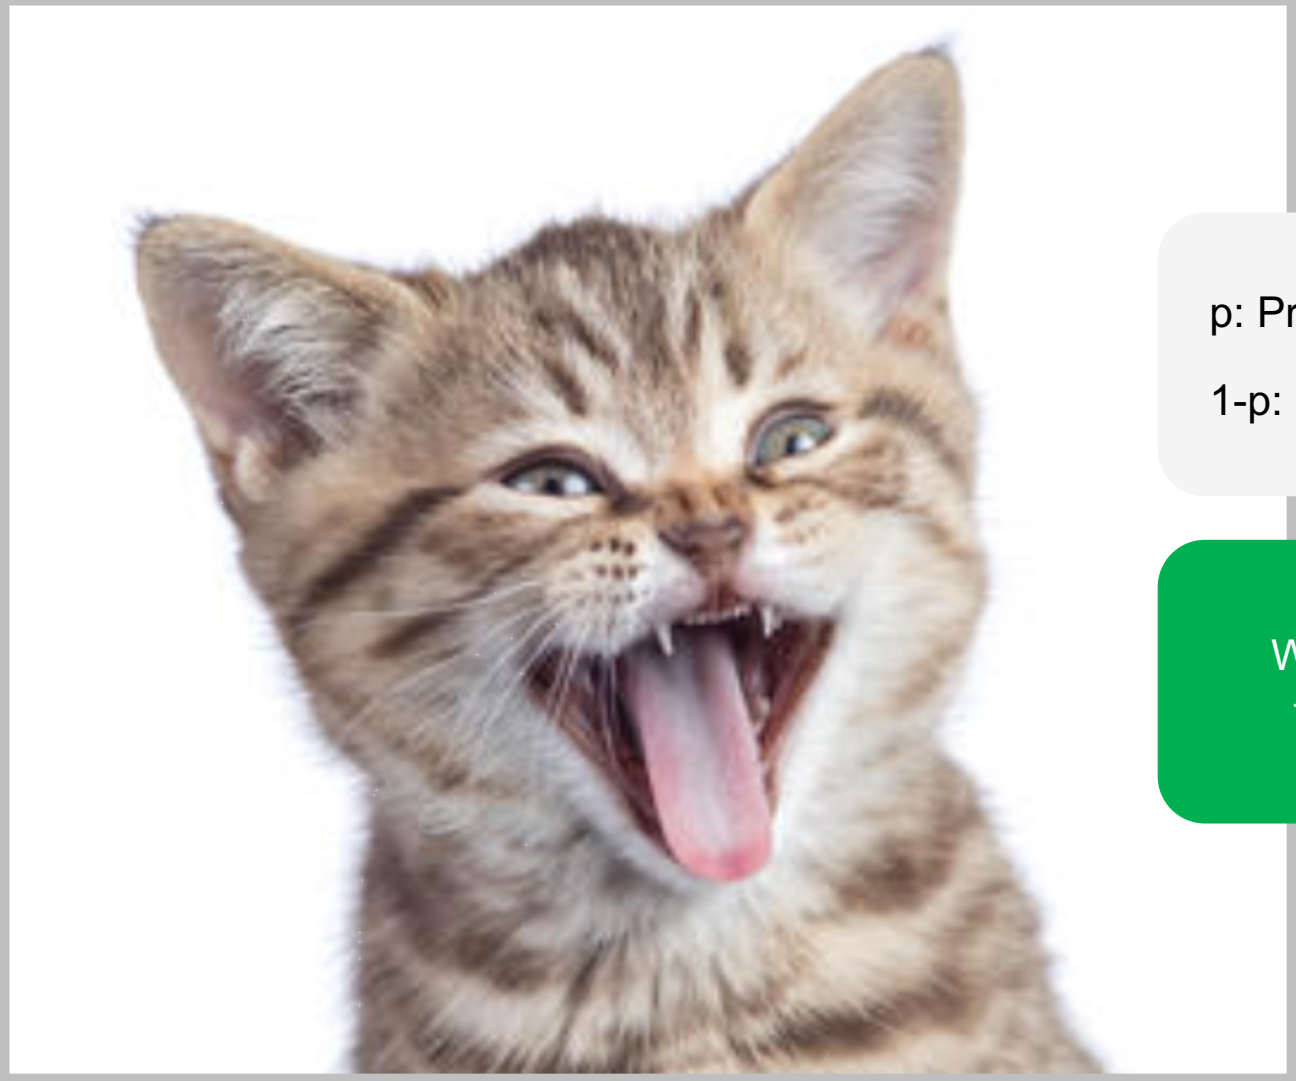

p: Prob. purring

1-p: Prob. grumpy

What is the best guess for the probability p?

How can I update my belief on p?

Prior distributions

The importance of priors in Bayesian Inference

$p(\boldsymbol{\mathcal{D}}|\boldsymbol{\theta})$: likelihood of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta})$: prior probability of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta}|\boldsymbol{\mathcal{D}})$: posterior of $\boldsymbol{\theta}$, given $\boldsymbol{\mathcal{D}}$

$p(\boldsymbol{\mathcal{D}})$: marginal probability of $\boldsymbol{\mathcal{D}}$

$p(y|x, \boldsymbol{\mathcal{D}})$: predictive distribution

Bayes' rule:

$$p(\theta|\boldsymbol{\mathcal{D}}) = \frac{p(\boldsymbol{\mathcal{D}}|\theta)\ p(\theta)}{P(\boldsymbol{\mathcal{D}})}$$

Prior distributions

The importance of priors in Bayesian Inference

$p(\boldsymbol{\mathcal{D}}|\boldsymbol{\theta})$: likelihood of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta})$: prior probability of $\boldsymbol{\theta}$

$p(\boldsymbol{\theta}|\boldsymbol{\mathcal{D}})$: posterior of $\boldsymbol{\theta}$, given $\boldsymbol{\mathcal{D}}$

$p(\boldsymbol{\mathcal{D}})$: marginal probability of $\boldsymbol{\mathcal{D}}$

$p(y|x, \boldsymbol{\mathcal{D}})$: predictive distribution

Bayes' rule:

$$p(\theta|\boldsymbol{\mathcal{D}}) = \frac{p(\boldsymbol{\mathcal{D}}|\theta)\ p(\theta)}{P(\boldsymbol{\mathcal{D}})}$$

A **prior probability distribution** of an uncertain quantity is the probability distribution that would express one's belief, before some evidence is taken into account.

→ For example, a prior could represent the distribution of votes coming from an opinion poll, prior to the election.

Priors: Subjective vs. Objective

Machine Learning Data Analytics

FAU

A **subjective prior** expresses the modeler's subjective belief.

- We formulate our (subjective) assumptions about modeling the data in terms of priors
- We have to work hard to understand the system under study in order to formulate our assumptions

An **objective prior** constrain prior beliefs to be "uninformative" about the parameters.

- The objective Bayes view is that formulating our assumptions is too difficult, especially in complex models

If we don't have strong beliefs about what θ should be, it is common to use an "uninformative" priors → **"Let the data speak for itself!"**

Priors: Informative vs. Uninformative

An **informative prior** expresses a specific information about a variable.

- For example, a reasonable informative prior about the temperature at noon tomorrow could be given by a normal distribution with expeced value equal to today's noon temperature and variance equal to the daily variance of the temperature.

An **uninformative prior** is designed to express vague or general information about a variable.

- For example, when tossing a coin, we assign the probability of 0.5 to both heads and tails.

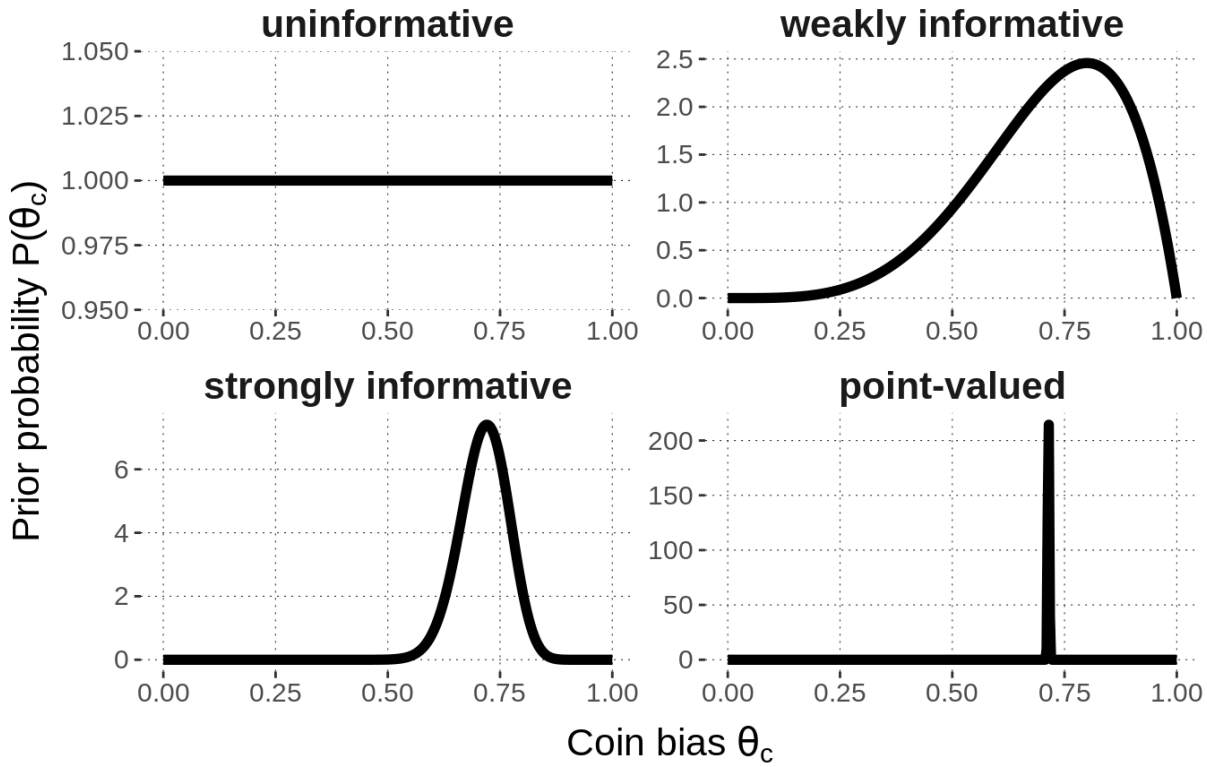

Image from: https://michael-franke.github.io/intro-data-analysis/Chap-03-03-models-parameters-priors.html

Hierarchical priors

The multi-level model

A key requirement for computing the posterior $p(\boldsymbol{\theta}|\boldsymbol{\mathcal{D}})$ is the specification of a prior $p(\boldsymbol{\theta}|\boldsymbol{\alpha})$, where alpha denotes a set of hyperparameters.

We have multiple levels of priors:

$$\alpha \rightarrow \theta \rightarrow \boldsymbol{\mathcal{D}}$$

$$p(\theta) = \int p(\theta|\alpha)\, p(\alpha)\, d\alpha$$

$$p(\boldsymbol{\mathcal{D}}) = \int p(\theta)\, p(\boldsymbol{\mathcal{D}}|\theta)\, \mathrm{d}\theta$$

The multi-level model

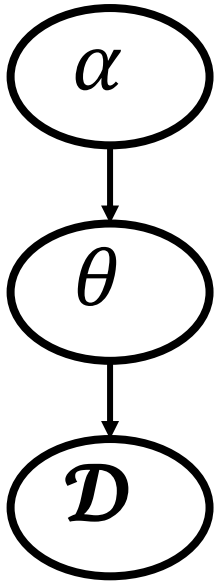

Hierarchical priors

An example: modeling cancer mortality in various cities

Consider the problem of predicting the cancer mortality rates in various cities. We measure the number of people N_i in various cities, as well as the number of people who died of cancer x_i in those cities. We assume that the mortality follows:

$$x_i \sim Bin(N_i, \theta_i).$$

A reasonable approach to estimate θ_i is that of assuming that they are drawn from the same distribution $\theta_i \sim Beta(a, b)$, where $\alpha = (a, b)$ are hyper-parameters in our model.

Then, the full joint distribution is written as

$$p(\boldsymbol{\mathcal{D}}, \theta, \alpha | N) = p(\alpha) \prod_{i=1}^{N} Bin(x_i | N_i, \theta_i) Beta(\theta_i | \alpha).$$

Note: It is crucial to infer α from the data itself.

Example taken from: Machine Leraning: A Probabilistic Perspective, Ch. 5.51

Empirical Prior

A computational “shortcut”

In hierarchical models, we need to **compute the posterior on multiple levels of variables**. For example,

$$p(\alpha, \theta|\boldsymbol{\mathcal{D}}) \propto p(\boldsymbol{\mathcal{D}}|\theta)\, p(\theta|\alpha)\, p(\alpha).$$

In some case, we can simplify the problems by **marginalizing over $\boldsymbol{\theta}$**. Then, we just need to compute:

$$p(\alpha|D).$$

As a computational “shortcut”, we can approximate the posterior on the hyper-parameters α with a point estimate. Since α is generally of a much smaller dimensionality than θ, it is less prone to overfitting and we safely assume a uniform prior on α.

$$\hat{\alpha} = \underset{\alpha}{\operatorname{argmax}}\, p(\boldsymbol{\mathcal{D}}, \alpha) = \underset{\alpha}{\operatorname{argmax}} \int p(\boldsymbol{\mathcal{D}}, \theta) p(\theta|\alpha) d\theta$$

Conjugate Prior

A computational “shortcut”

A prior $p(\theta)$ is a **conjugate prior** for a particular likelihood $p(y|\theta)$ if the resulting posterior $p(\theta|y)$ has the same algebraic form.

Conjugate priors are widely used because they provide advantages:

- they usually allow us to derive a closed-form expression for the posterior distribution;
- they are easy to interpret,

Note: Conjugate priors simplify the computation, but are often not flexible enough to encode our prior knowledge → We can also use mixture of conjugate priors.

Conjugate Prior

An example

Likelihood (Binomial): $p(\mathcal{D}|\theta) = Bin(\theta) = \binom{n}{x}\theta^x(1-\theta)^{n-x}$

Prior (Beta): $p(\theta) = Beta(\alpha,\beta) = \frac{1}{B(\alpha,\beta)}\theta^{\alpha-1}(1-\theta)^{\beta-1}$

We plug them into the Bayes' formula to derive the posterior distribution:

$$p(\theta|\mathcal{D}) = \frac{p(\mathcal{D}|\theta)\,p(\theta)}{\int p(\mathcal{D}|\theta)\;p(\theta)d\theta}$$

$$= \frac{\binom{n}{x}\theta^x(1-\theta)^{n-x}\;\frac{1}{B(\alpha,\beta)}\theta^{\alpha-1}(1-\theta)^{\beta-1}}{\int\binom{n}{x}\theta^x(1-\theta)^{n-x}\;\frac{1}{B(\alpha,\beta)}\theta^{\alpha-1}(1-\theta)^{\beta-1}\;d\theta}$$

$$= \frac{\frac{n\,C_x}{B(\alpha,\beta)}\;\theta^{x+\alpha-1}\,(1-\theta)^{n-x+\beta-1}}{\frac{n\,C_x}{B(\alpha,\beta)}\;\int\theta^{x+\alpha-1}\,(1-\theta)^{n-x+\beta-1}\;d\theta} = Beta(x+\alpha, n-x+\beta)$$

Prior: Beta(2, 2)

Prior: Beta(2, 2)

Posterior: Beta(2+2, 4+2) = Beta(4, 6)

Prior: Beta(2, 2)

Posterior: Beta(2+2, 4+2) = Beta(4, 6)

Example inspired by: https://towardsdatascience.com/understanding-conjugate-priors-21b2824cddae

Bayesian Inference
Linear Regression (Bayesian treatment)

Linear Regression (Bayesian treatment)

Given the observed data $\mathcal{D} = \{x^{(n)}, y^{(n)}\}$, we assume to know the noise variance σ^2.

We would like to compute the posterior over the parameters, i.e,

$$p(w|\mathcal{D}, \sigma^2).$$

(We assume throughout a Gaussian likelihood model).

In linear regression **the likelihood is given by:**

$$p(y|X, w, \mu, \sigma^2) = \mathcal{N}(y|\mu + Xw,\ \sigma^2 I_N)$$

$$\propto \exp(-\frac{1}{2\sigma^2}(y - \mu - Xw)^T(y - \mu - Xw))$$

where μ is an offset term.

Linear Regression (Bayesian treatment)

The conjugate prior of a Gaussian likelihood is also Gaussian*, which we will denote by

$$p(w) = \mathcal{N}(w|w_0, V_0).$$

Using the Bayes rule for Gaussian*, the posterior is given by

$$p(w|X, y, \sigma^2) \propto \mathcal{N}(w|w_0, V_0)\, \mathcal{N}(\mathrm{y}|\mathrm{Xw}, \sigma^2 I_N) = \mathcal{N}(w|w_N, V_N)$$

where

$$w_N = V_N V_0^{-1} w_0 + \frac{1}{\sigma^2} V_N X^T y$$

$$V_N = \sigma^2 (\sigma^2 V_0^{-1} + X^T X)^{-1}$$

* See: Murphy K., „Machine Learning: A Probabilistic Perspective“ (2012)

Linear Regression (Bayesian treatment)

The posterior predictive distribution at a test point x is given by *

$$p(y|x, \boldsymbol{\mathcal{D}}, \sigma^2) = \int \mathcal{N}(y|x^T w, \sigma^2)\mathcal{N}(w|w_N, V_N)dw$$

$$= \mathcal{N}(y|w_N^T x, \sigma_N^2(x))$$

where $\sigma_N^2(x) = \sigma^2 + x^T V_N x$.

The variance in this prediction depends on the variance of the observation noise, σ^2, and the variance in the parameters, $\mathrm{V_N}$.

Bayesian Inference
Recap

Recap

- Bayesian modelling
  - Prior
  - Posterior
  - Likelihood
- Bayesian modelling requires integration over parameters
  - For complex models it could be not tractable! (We cannot compute the integral analytically)
- Priors
  - Informative vs Uninformative
  - Conjugate priors
- Linear regression with Bayesian treatment