



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Lecture Pattern Analysis

Part 01: Vocabulary, Probabilities, and Sampling

Christian Riess

IT Security Infrastructures Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 28, 2022



Introduction

- We require some vocabulary for our communication about Machine Learning (ML)
- We also require a formal framework to discuss ML algorithms on a scientific basis
- Arguably the most widely used framework is probability theory
- We will also introduce our first algorithm, namely how to sample from a Probability Density Function (PDF)

Pattern Recognition Recap and Classification Vocabulary

- Remember the steps of the classical pattern recognition pipeline:



- Fundamental ML assumption: good feature representations map similar objects to similar features
- Classifier training is virtually always **supervised**, i.e. a training sample is a tuple (\mathbf{x}_i, y_i) (cf. lecture “Pattern Recognition”)
- Unsupervised** ML works without labels, i.e., it only operates on inputs (\mathbf{x}_i) . Hence, unsupervised ML only works on the distribution of the features
- Fashionable variants are semi-supervised ML (some data has labels), self-supervised ML (auto-generate surrogate labels)

Recap on Probability Vocabulary

- We oftentimes operate with random variables X, Y
- Important vocabulary and equations are:

Joint distribution $p(X, Y)$

Conditional distribution of X given Y $p(X|Y)$

Sum rule / marginalization over Y $p(X) = \sum_Y p(X, Y)$

Product rule $p(X, Y) = p(Y|X) \cdot p(X)$


Bayes rule $p(Y|X) = \frac{p(X|Y) \cdot p(Y)}{p(X)}$

Bayes rule in the language of ML $\text{posterior} = \frac{\text{likelihood} \cdot \text{prior}}{\text{evidence}}$

- Please browse the book by Bishop, Sec. 1.2.3, to refresh your mind if necessary!

Sampling from a PDF

- Oftentimes, it is necessary to draw samples from a PDF
- Example:
 - Logistic Regression fits a single regression curve to the data (cf. PR)
 - Bayesian Logistic Regression fits a distribution of curves



The distribution is narrow at observations (crosses), and wider otherwise



- Sample curves from the distribution to obtain its spread (“uncertainty”)
- Special PDFs like Gaussians have closed-form solutions for sampling
- We look now at a sampling method that works on **arbitrary PDFs**

Idea of the Sampling Algorithm

- The key idea is to use the cumulative density function (CDF) $P(z)$ of $p(X)$,

$$P(z) = \int_{-\infty}^z p(X) dX \quad (1)$$

- A sample uniformly drawn from the CDF y -axis intersects $P(z)$ at location z
- This z position is our random draw from $p(x)$:



Sampling Algorithm

- Discretize the domain of the PDF $p(X)$
- Linearize $p(X)$ if it is multivariate
- Calculate the cumulative density function $P(z)$ of $p(X)$, the range of that CDF must be between 0 to 1
- Draw a uniformly distributed number u between 0 and 1
- The sample from the PDF is

$$z^* = \underset{z}{\operatorname{argmin}} u \leq P(z) \quad (2)$$

- Additional reference / strictly optional:
Bishop Chapter 11 contains many more (also advanced) sampling strategies



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Lecture Pattern Analysis

Part 02: Non-Parametric Density Estimation

Christian Riess

IT Security Infrastructures Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 28, 2022



Introduction

- Density Estimation = create a PDF from a set of samples
- The lecture Pattern Recognition introduces parametric density estimation:
 - Here, a parametric model (e.g., a Gaussian) is fitted to the data
 - Maximum Likelihood (ML) estimator:

$$\theta^* = \underset{\theta}{\operatorname{argmax}} p(\mathbf{x}_1, \dots, \mathbf{x}_N | \theta) \quad (1)$$


- Maximum a Posteriori (MAP) estimator:

$$\theta^* = \underset{\theta}{\operatorname{argmax}} p(\theta | \mathbf{x}_1, \dots, \mathbf{x}_N) \stackrel{\text{Bayes}}{=} \frac{p(\mathbf{x}_1, \dots, \mathbf{x}_N | \theta) \cdot p(\theta)}{p(\mathbf{x}_1, \dots, \mathbf{x}_N)} \quad (2)$$

- Browse the PR slides if you like to know more
- Parametric density estimators require a good function representation
- Non-parametric density estimators can operate on arbitrary distributions

Non-Parametric Density Estimation: Histograms

- Non-parametric estimators do not use functions with a limited set of parameters
- A **simple non-parametric baseline** is to create a histogram of samples¹
 - The number of bins is important to obtain a good fit



- **Pro:** Good for a quick visualization
- **Pro:** “Cheap” for many samples in low-dimensional space
- **Con:** Discontinuities at bin boundaries
- **Con:** Scales poorly to high dimensions (cf. curse of dimensionality later)

¹ See introduction of Bishop Sec. 2.5

Improving on the Histogram Approach

- A kernel-based method and a nearest-neighbor method are slightly better
- Both variants share their mathematical framework:

- Let $p(\mathbf{x})$ be a PDF in D -dim. space, and R a small region around \mathbf{x}
 \rightarrow The probability mass in R is $p = \int_R p(\mathbf{x}) d\mathbf{x}$
- Assumption 1: in R are many points $\rightarrow p$ is a relative frequency,

$$p = \frac{\text{\# points in } R}{\text{total \# of points}} = \frac{K}{N} \quad (3)$$

- Assumption 2: R is small enough s.t. $p(\mathbf{x})$ is approximately constant,

$$p = \int_R p(\mathbf{x}) d\mathbf{x} = p(\mathbf{x}) \int_R d\mathbf{x} = p(\mathbf{x}) \cdot V \quad (4)$$

- Both assumptions together are slightly contradictory, but they yield

$$p(\mathbf{x}) = \frac{K}{N \cdot V} = \frac{\text{\# points in } R}{\text{total \# of points} \cdot \text{Volume of } R} \quad (5)$$

Kernel-based DE: Parzen Window Estimator (1/2)

- The Parzen window estimator fixes V and leaves K/N variable²
- D -dimensional Parzen window kernel function (a.k.a. “box kernel”):

$$k(\mathbf{u}) = \begin{cases} 1 & \text{if } |u_i| \leq \frac{1}{2} \quad \forall i = 1, \dots, D \\ 0 & \text{otherwise} \end{cases} \quad (6)$$

- Calculate K with this kernel function:

$$K(\mathbf{x}) = \sum_{i=1}^N k\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right) \quad (7)$$

where h is a scaling factor that adjusts the box size

- Hence, the whole density is

$$p(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \frac{1}{h^D} k\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right) \quad (8)$$

²See Bishop Sec. 2.5.1

Kernel-based DE: Parzen Window Estimator (2/2)

- The kernel removes much of the discretization error of the fixed-distance histogram bins, but it still leads to blocky estimates
- Replacing the box kernel by a Gauss kernel further smooths the result,

$$p(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \frac{1}{2\pi h^2} \cdot \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}_i\|_2^2}{2h^2}\right), \quad (9)$$

where h^2 is the standard deviation of the Gaussian

- Mathematically, also any other kernel is possible if these conditions hold:

$$k(\mathbf{u}) \geq 0 \quad (10)$$

$$\int k(\mathbf{u}) d\mathbf{u} = 1 \quad (11)$$

K-Nearest Neighbors (k-NN) Density Estimation

- Recall our derived equation for estimating the density

$$p(\mathbf{x}) = \frac{K}{N \cdot V} = \frac{\text{\# points in } R}{\text{total \# of points} \cdot \text{Volume of } R} \quad (12)$$

- The Parzen window estimator fixes V , and K varies
- The k-Nearest Neighbors estimator fixes K , and V varies
- k-NN calculates V from the distance of the K nearest neighbors³
- Note that both the Parzen window estimator and the k-NN estimator are “non-parametric”, but they are not free of parameters
 - The kernel scaling h and the number of neighbors k are **hyper-parameters**, i.e., some form of prior knowledge to guide the model creation
 - The model parameters are the samples themselves. Both estimators need to store all samples, which is why they are also called **memory methods**

³See Bishop Sec. 5.2.2

First Glance at the Model Selection Problem

- Optimizing the hyperparameters is also called **Model Selection Problem**
- Supervised methods use cross validation (CV) via Maximum Likelihood (ML)
- We can use CV to optimize the DE hyperparameters by using the **prediction of held-out samples** as objective function:

- Split the data into J folds:

$$S_{\text{train}}^j = S \setminus \{x_{\lfloor \frac{N}{J} \rfloor \cdot j}, \dots, x_{\lfloor \frac{N}{J} \rfloor \cdot (j+1) - 1}\},$$

$$S_{\text{test}}^j = S \setminus S_{\text{train}}^j$$

- Let α be the unknown hyperparameters, and let $p_j(\mathbf{x}|\alpha)$ be the density estimate for samples S_{train}^j on hyperparams α
- Then, the ML estimate is

$$\alpha^* = \underset{\alpha}{\operatorname{argmax}} \prod_{j=1}^J \prod_{\mathbf{x} \in S_{\text{test}}^j} p_j(\mathbf{x}|\alpha) \quad (13)$$

- In practice, take the logarithm (“log likelihood”) to mitigate numerical issues
 → the product becomes a sum



FRIEDRICH-ALEXANDER-
UNIVERSITÄT
ERLANGEN-NÜRNBERG
SCHOOL OF ENGINEERING

Lecture Pattern Analysis

Part 03: Bias and Variance

Christian Riess

IT Security Infrastructures Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg

April 28, 2022



Introduction

- We search in the exercise for good kernel parameters via cross validation
- This model selection task touches the question of **generalization** to new data:
 - Too large kernel: smears out the structure of the training data, but covers new samples
 - Too small kernel: closely follows the training data, but might miss new samples
 - “Right” kernel size: represents the structure of the training data and also covers new samples
(to the extent possible with the given method and data)
- This is an instance of the **bias-variance tradeoff**¹

¹ See PR lecture or Hastie/Tibshirani/Friedman Sec. 7-7.3 if more details are desired


Bias and Variance in Regression

- **Bias** is the square of the *average deviation* of an estimator from the ground truth
- **Variance** denotes is the *variance of the estimates*, i.e., the expected squared deviation from the estimated mean²
- Informal interpretation:
 - High bias indicates **model undercomplexity**: we obtain a poor fit to the data
 - High variance indicates **model overcomplexity**: the fit also models not just the structure of the data, but also its noise
- Higher model complexity (= more model parameters) tends to lower bias and higher variance
- We will usually not be able to get bias and variance simultaneously to 0
- Regularization increases bias and lowers variance


²See Hastie/Tibshirani/Friedman Sec. 7.3 Eqn. (7.9) for a detailed derivation

Sketches for Model Undercomplexity and Overcomplexity

Ground truth




Underfitting



Measurements



Overfitting



- Note that this example implicitly contains a smoothness assumption
- It does not claim that there is a universally best fit on arbitrary input distributions (because of the No-Free-Lunch Theorem)

Transferring Bias and Variance to our Density Estimators

- Our kernel framework can directly replicate these investigations by retargeting our kernels to regression or classification:
- Regression:
 - Estimate $f(\mathbf{x})$ at position \mathbf{x} as a kernel-weighted sum of the neighbors or
 - as a k -NN mean of k neighbors
- Classification:
 - Estimate for classes c_1 and c_2 individual densities, evaluate $p_{c_1}(\mathbf{x})$ and $p_{c_2}(\mathbf{x})$, and select the class with higher probability or
 - Select the majority class within k nearest neighbors
- We will then observe that
 - Larger kernel support / larger k increases bias and lowers variance
 - Smaller kernel support / smaller k lowers bias and increases variance
- Analogously, we can use the notion of bias/variance also in our exercise task of unsupervised density estimation

Bonus Slide: Current Research

- Just if you are curious — this is **strictly voluntary**:
- Here is a video on the recent ICML paper “Maximum Likelihood With Bias-Corrected Calibration is Hard-To-Beat at Label Shift Adaptation” by Avanti Shrikumar from Stanford University:
<https://www.youtube.com/watch?v=ZBXjE9QTruE>
- It adapts the prior of a deep learning model without retraining the model, using tools from our lecture
- I find the talk also generally instructive as an academic presentation
- Enjoy!